



GAME PLAYING

林伯慎 Bor-shen Lin

bslin@cs.ntust.edu.tw

<http://www.cs.ntust.edu.tw/~bslin>

AGENDA

- Game Trees
- Assumptions
- Static evaluation functions
- Searching game trees
- Minimax
- Bounded lookahead
- Alpha-beta pruning
- Checkers




GAME TREES

- Game trees are used to represent two-player games.
- Alternate moves in the game are represented by alternate levels in the tree (plies).
- Nodes in the tree represent positions.
- Edges between nodes represent moves.
- Leaf nodes represent won, lost or drawn positions.
- For most games, the game tree can be enormous.



GAME TREES




This is an example of a partial game tree for the game tic-tac-toe. Even for this simple game, the game tree is very large.

ASSUMPTIONS

- The opponent is **rational & adversarial** – will play to win.
- The game is **zero-sum** – if one player wins, the other loses.
- Usually, the two players **have complete knowledge of the game**. For games such as poker, this is clearly not true due to **element of chance**.



NON-ADVERSARIAL OPPONENT



- Play golf with the boss or customers
- Play cards or chesses with children



STATIC EVALUATION FUNCTIONS

- Since game trees are too large to be fully searched, it is important to have a function to statically evaluate a given position in the game.
- Are usually **weighted linear function**
 - Can be extremely complex
- A static evaluator assigns a score to a position:
 - High positive = computer is winning
 - Zero = even game
 - High negative = opponent is winning



EVALUATION FUNCTION

- A static evaluator
 - give a better score to a better position – the actual values are not so important.
- How ?
 - could be assigned manually based on expert knowledge
 - could be learned according to the result of the game.
- Chess
 - q = number of queens
r = number of rooks
n = number of knights
b = number of bishops
p = number of pawns
score = $9q + 5r + 3b + 3n + p$



SEARCH OF GAME TREES

- The concept of **ply**: the depth of search tree
 - Look ahead more plies → higher depth → higher complexity
- Huge search space
 - Branching factor is usually high
 - Exhaustively searching a game tree is not a good idea.
 - Even for a game as simple as tic-tac-toe there are over 350,000 nodes ($9 \cdot 8 \cdot \dots \cdot 1$) in the complete game tree .
- Difference of state space search between *game playing* and *problem solving*
 - Computer only gets to choose **every other move** through the tree – the opponent chooses the others.




MINIMAX ALGORITHM

- Minimax is used to evaluate game trees.
 - Already applied in risk management of investment
 - MAX: Computer's turns to maximize the score
 - MIN: Opponent's turns to minimize the score
 - The opponent is **not cooperative** to make the computer win, but adversarial to defeat the computer
- A static evaluator
 - is applied to **leaf nodes**, and values are passed back up the tree to determine the best score the computer can obtain against a rational opponent.
- The game tree is searched in **depth-first** way.
 - With a predetermined depth



EXAMPLE OF MINIMAX



MINIMAX ALGORITHM

```
Function minimax(current_node)
{
    if is_leaf(current_node)
        then return static_evaluation(current_node);
    if is_min_node(current_node)
        then return min(minimax(children(current_node)));
    if is_max_node(current_node)
        then return max(minimax(children(current_node)));
}
```



BOUNDED LOOKAHEAD

- For trees with **high depth** or **very high branching factor**, minimax cannot be applied to the entire tree.
- In such cases, bounded lookahead is applied
 - When search reaches **a specified depth**, the search is **cut off**, and the static evaluator applied.
(Not search to leaf nodes)
- Corresponds to human's way
 - The masters consider only important states instead of checking all possible states exhaustively




BOUNDED MINIMAX ALGORITHM

```
Function minimax(current_node)
{
    if is_leaf(current_node)
        then return static_evaluation(current_node);
    if depth_of(current_node) == max_depth
        then return static_evaluation(current_node);
    if is_min_node(current_node)
        then return min(minimax(children(current_node)));
    if is_max_node(current_node)
        then return max(minimax(children(current_node)));
}
```



CUT-OFF AT FIXED DEPTH

- Not good to cut-off sometimes
 - Search deeper states
- Solution
 - Not fixed depth, but till **quiescent positions**
 - Quiescent position: the next move is unlikely to cause a large change in the relative positions of the two players.



ALPHA-BETA PRUNING

- A method that can often cut off a large part of the game tree.
- Based on the idea that if a move is **clearly bad**, there is no need to waste time to search the consequences of it completely.



PROCESS OF FINDING MAXIMUM

- *alpha = -10000000; // very small value*
argmax = 0;
for(int i = 0; i < X.length; i++) {
 if(X[i] > alpha) { // replace alpha if larger
 alpha = X[i]; // replaced by X[i]
 argmax = i; // storing its index
 }
}
- *alpha* is the *temporary maximum* during the execution.
- *alpha* can *increase only* (could be replaced by a larger entry) during the the maximization process.
- *alpha* is the maximum *finally* after all elements are visited



EXAMPLE

→
X [6, 15, 9, 23, 8, 37, 19]
|
6 > $-\infty$?
 α : $-\infty$

- Value of α is updated when larger values are tested (compared)
- Status of α
 - $-\infty, 6, 15, 23, 37$
 - keep increasing
- α is the temporary maximum



FINDING MINIMUM IN AN ARRAY

- *beta = 10000000; // very large value*

argmin = 0;

for(int i = 0; i < X.length; i++) {

if(X[i] < beta) { // replace it if smaller than beta

beta = X[i]; // replaced by X[i]

argmin = i; // storing its index

}


}

- *beta* is the *temporary minimum* during the execution.
- *beta* can *decrease only*
- *beta* is the maximum *finally* after all elements are visited




ALPHA-BETA PRUNING (CONT'D)

- Recursive minmax algorithm is depth-first search
- DFS on game tree **optimized** with alpha-pruning
- Good player does not exhaustively or blindly search the game tree.
 - Pruning & optimized tree search




ALPHA-BETA PRUNING



- Having examined the nodes with values 7 and 1, there is no need to examine the final node.
- MIN node
 - $\beta=1 \leq \alpha=3$




ALPHA-BETA PRUNING (CONT'D)



- MAX nodes $\rightarrow \alpha$, MIN nodes $\rightarrow \beta$
- Initial values : α -Infinity, β Infinity
- Pruning check
 - MIN nodes: $\beta < \text{parent's } \alpha$
 - MAX nodes: $\alpha \geq \text{parent's } \beta$



ALPHA-BETA PRUNING (CONT'D)



EFFECTIVENESS OF ALPHA-BETA PRUNING

- Ideal case

- First branch can be pruned
- Reduce the branching factor from b to $b^{1/2}$.
- Double the search depth under the same computational cost

- Deep blue

- Reduce the branching factor from 38 to 6 by applying alpha-beta pruning



REDUCING SEARCH COMPLEXITY

- Complexity of search tree
 - Determined by **depth** and **breadth** (b^d)
- Reduction of complexity
 - Control of depth: force the search to stop at certain level
 - Control of breadth: limit the number of choices when spanning the tree



GAMES OF CHANCE

- The methods described so far do not work well with games of chance such as poker or backgammon.
- Expectminimax is a variant of minimax designed to deal with chance.
- Nodes have expected values based on probabilities.



CHECKERS

- In 1959, Arthur Samuel published a paper in which he described a computer program that could play checkers to a high level using **minimax** and **alpha-beta pruning**.
- Chinook, developed in Canada defeated the world champion:
 - Uses alpha-beta pruning.
 - Has a database of millions of end games.
 - Also has a database of openings.
 - Uses heuristics and knowledge about the game.



CHECKERS

○ Blondie 24

- Use evolutionary techniques to develop neural networks that can learn how to play checkers and how to win
- Use standard Minimax algorithm
- The static evaluation function is evaluated by neural networks (machine learning)
- Is able to beat many human players, but not nearly at the level of Chinook
- Without any help from humans



CHESS

- In 1997, Deep Blue defeated world champion, Garry Kasparov.
- This has not yet been repeated.
- Current systems use parallel search, alpha-beta pruning, databases of openings and heuristics.
- The deeper in a search tree the computer can search, the better it plays.



GO (WEI-QI)

- Ancient Oriental game
- Complex game played on a 19 x19 board.
- Average branching factor in search tree around 360 (compared to 38 for chess).
- Potentially Infinite depth
 - Ko(打劫): infinite loop occurs
- Methods use pattern matching or selective search to explore the most appropriate parts of the search tree.
- Due to vast complexity, traditional “brute force” methods don’t work on Go
- The Everest of search complexity
 - A number greater than there are atoms in the universe



GO (CONT'D)

- Goals are vague and complicated
- Not easy to define *evaluation function*
 - Almost the same number of stones for black & white
 - Difficult to define *threatens*
- Not easy to judge the strengths/weakness (the state of game)
 - Power (~ investments): influence, potential for earning profits
 - Profits (~assets)



SUMMARY

- Game tree can be used to represent two-player games.
- Searching game trees is very hard for all but the simplest games.
- Minimax is an algorithm that identifies the best move to make in a two-player game with perfect knowledge, assuming the entire tree can be examined.
- When the entire tree cannot be examined, a static evaluator is needed that can assign a score to any given position according to how well each player is doing and how likely each player is to win the game from that position.



SUMMARY

- Alpha-beta pruning enables Minimax to run more efficiently because it removes unnecessary branches from the game tree.
- In the best case, alpha-beta pruning enables Minimax to search the game tree to double the depth that it would be able to search without pruning.



SUMMARY

- Games such as Go-moku and ti-tac-toe have been *solved*, meaning that the result of a game is known from the start.
- Computers are able to beat the best players in the world at games such as Othello and checkers.
- Although Deep Blue beat Garry Kasparov in 1997, computers can win any professional human players.
- *AlphaGo* defeat *Lee Sedol* (9d professional player) on go game on *2016/3/9*.



MONTE-CARLO METHOD

- Repeated *random sampling* to obtain numerical results
- Useful for those problems difficult or impossible to solve using other mathematical methods
- Used in optimization, numerical integration, probability distribution



Monte-Carlo Tree Search

- Converged slowly to *minimax*
- Can *minimize the search space* like alpha-beta pruning
- Do not need explicit evaluation function (game-end condition) as computer game of chess
- Concentrates on *most promising subtrees*
- Can be interrupted at any time and yields most promising move already found
- Successfully applied to computer-go since 2006



ALPHA-GO

- Tree search guided by two deep neural networks
 - Millions of neurons
- Policy network
 - Goal: narrow the search to consider only the moves most likely to lead to a win
 - Trained by 30 millions moves
 - Discover new strategies by self playing games
- Value network
 - Goal : estimate the winner in each position
- Trained by human and computer play
 - Allow it to improve itself by just watching the playing games
 - Reinforcement learning from games of self play




ALPHA-GO

- Google Cloud Platform

- Distributed & parallel processing with multiple CPUs/GPUs
- Open source library
- Data flow graphs build with Python (TensorFlow)
- tensorflow.org



ALPHA-GO



Cited from *Google*

<http://googleresearch.blogspot.tw/2016/01/alphago-mastering-ancient-game-of-go.html>

- In the past, computers can defeat only amateurs
- AlphaGo defeated almost the smartest human expert



FURTHER THINKING

- Applied to business competition or investment?
 - Your opponents hope to win
 - How to represent the state
 - How to evaluate the state
 - Search and make decision automatically
- How is *capability* defined in game?
 - Capabilities of players are measured through games
 - Probability to win → capability moves up/down
 - Another way of *measuring intelligence* !

