



FUNDAMENTALS OF ARTIFICIAL NEURAL NETWORK

林伯慎 Bor-shen Lin

bslin@cs.ntust.edu.tw

<http://www.cs.ntust.edu.tw/~bslin>

ARTIFICIAL NEURAL NETWORK, WHY?


- Mimic how **human brains** function
 - Human brain is superior to digital computers in many ways
 - A baby can recognize many objects
 - Computers are good at arithmetic
- Human brains are
 - **Robust** and **fault-tolerant**
 - **Flexible** and **adaptive** (learning)
 - Can deal with fuzzy, probabilistic, noisy or inconsistent information
 - Small, compact & consumes low power



HUMAN BRAIN

- 50 to 150 billion neurons in brain
- Neurons grouped into networks
 - Axons send outputs to cells
 - Received by dendrites, across synapses

Figure 12.4 Portion of a Network: Two Interconnected Biological Cells



FUNCTIONS OF NEURONS

○ Transmission of Signal

- Complicated chemical process
- Transmitter **substances** are released from the sending site
- Raise or lower **electrical potential** inside the body of receiving cell
- If the potential reaches a threshold, a **pulse** of a fixed strength and duration is sent (**fired**)
- After firing, the cell has to wait a time called **refractory period** before it can fire again



MATHEMATICAL MODEL


- $y_i(t+\Delta t) = g(\sum_j w_{ji}x_j(t) - \mu_i)$.
- Activation function $g(h)$
$$g(h) = \begin{cases} 1 & \text{if } h \geq 0 \\ 0 & \text{otherwise.} \end{cases}$$
- w_{ji} : the strength of synapse connecting neuron j to neuron i
- μ_i : threshold value for neuron i
- If the weighted sum of the inputs to the neuron i exceeds the threshold μ_i , the neuron fires and the output becomes 1.




PERCEPTRON: ARTIFICIAL NEURON

- Artificial neurons are based on biological neurons.
- Each neuron in the network receives one or more inputs.
- An activation function is applied to the inputs, which determines the output of the neuron – the activation level.


- The charts on the right show three typical activation functions $g(\cdot)$.



(a) Step function



(b) Sigmoid function



(c) Linear function

ACTIVATION FUNCTION

○ Characteristics

- Monotonic (with positive slope)
- Range: between 0 and 1
- Differentiable s.t. easier for optimization

○ Popular types

- Sigmoid

$$g(h) = \frac{1}{1+e^{-h}}$$

- Hyperbolic tangent (denoted as \tanh)

$$g(h) = \frac{e^h - e^{-h}}{e^h + e^{-h}}.$$


- Rectified linear (denoted as RELU)

$$g(h) = \max(0, h)$$



PERCEPTRON

- A perceptron is a single neuron that classifies a set of inputs into one of two categories.
- If the inputs are in the form of a grid, a perceptron can be used to **recognize visual images of shapes**.
- The perceptron usually uses a step function, which returns 1 if the weighted sum of inputs exceeds a threshold, and -1 otherwise.



LEARNING OR FUNCTION

○ OR

- $0\ 0 \rightarrow 0$
- $0\ 1 \rightarrow 1$
- $1\ 0 \rightarrow 1$
- $1\ 1 \rightarrow 1$

○ Initial setting

- $w_1 = -0.2$
 $w_2 = 0.4$
 $a = 0.2$


Epoch	X1	X2	Expected Y	Actual Y	Error	w1	w2
1	0	0	0	0	0	-0.2	0.4
1	0	1	1	1	0	-0.2	0.4
1	1	0	1	0	1	0	0.4
1	1	1	1	1	0	0	0.4
2	0	0	0	0	0	0	0.4
2	0	1	1	1	0	0	0.4
2	1	0	1	0	1	0.2	0.4
2	1	1	1	1	0	0.2	0.4
3	0	0	0	0	0	0.2	0.4
3	0	1	1	1	0	0.2	0.4
3	1	0	1	1	0	0.2	0.4
3	1	1	1	1	0	0.2	0.4

TRAINING A PERCEPTRON

- A perceptron can be trained as follows:
 - First, inputs are given random weights (usually between -0.5 and 0.5).
 - An item of training data is presented. If the perceptron mis-classifies it, the weights are modified according to the following:
$$w_i' \leftarrow w_i + a \cdot x_i \cdot e \text{ (} e = d - y, d : \text{desired output)}$$
 - e is the size of the error
 - a is the learning rate, between 0 and 1 .
- $h_i' = w_i' x_i = (w_i + ax_i e)x_i = w_i x_i + ax_i^2 e = h_i + ax_i^2 e$
if $e < 0 : d < y, y$ too big \rightarrow updated weight w_i'
 \rightarrow decrease h ($h_i' < h_i$) \rightarrow decrease y (since $g(\cdot)$ monotonic)



LINEAR DECISION FUNCTION OF A PERCEPTRON




- A family of linear functions with adjustable parameters w_1 , w_2 and μ .
- A learning procedure can be performed to update the parameters of the linear function so as to find out the decision boundary that separates the 2D points of the two classes.




DECISION FUNCTIONS FOR CLASSIFIERS

Decision boundary $f(x, y) = 0$
can separate the 2-D plane into
two regions (+ or -).




$$f(x, y) = ax + by + c$$

$(a, b > 0)$



$$f(x, y) = (x - x_0)^2 + (y - y_0)^2 - r^2$$

$$f(x_0, y_0) = -r^2 < 0$$





$$f(x, y) = y - ax^2 - bx - c$$

$$(b^2 - 4ac > 0)$$





FAMILY OF FUNCTIONS



- Rotate the line by changing w_1 and w_2
- Shift the line by changing b




HOW DOES A NEURON LEARN?



- $w_1x_1 + w_2x_2 + b = 0$
- Change orientation by learning w_1 and w_2
- Shift location by learning b



PROBLEM OF LINEAR CLASSIFIER





- NOT linearly separable (for single neuron)



LIMITATION OF A PERCEPTRON

- A Single Perceptron can only classify **linearly separable** functions using step function.
- Can **NOT** model the problem that is not linearly separable such as simple Exclusive-OR.
 - **No line** can separate the black/white dots or EOR



$$X = \sum_i w_i \cdot x_i$$

$$Y = +1 \text{ if } X > t$$


$$- 1 \text{ if } X < t$$

$$\sum_i w_i \cdot x_i > t \quad +1$$

$$\sum_i w_i \cdot x_i < t \quad -1$$



MULTIPLE LAYERS OF NEURONS




- Multiple neurons (every one is linear)
- Multiple layers to form decision boundary



MULTILAYER PERCEPTRONS (MLP)


- Multilayer neural networks can classify a range of functions, including nonlinearly separable ones.
- Each input layer neuron connects to all neurons in the hidden layer.
- The neurons in the hidden layer connect to all neurons in the output layer.




A feed-forward network

GRADIENT DESCENT (FOR MINIMIZATION)

- To find minimum of $f(x)$ from arbitrary point x
 - **Local minimum** can be obtained



If $f'(x) < 0$, increase x ($dx > 0$)




If $f'(x) > 0$, decrease x ($dx < 0$)

$\Delta x = -\epsilon(df/dx)$ to find minimal $f(x)$

$x_{new} = x_{old} + \Delta x, \Delta x = -\epsilon(df/dx)$



MEANING OF DERIVATIVE



- Try to increase scores by studying more
- Higher derivative: $r'(s) = \frac{dr}{ds}$ large
 - $dr = r'(s)ds \rightarrow ds$ leads to more increase of r (higher dr)
- Lower derivative : $r'(s) = \frac{dr}{ds}$ small
 - $dr = r'(s)ds \rightarrow ds$ leads to less increase of r (lower dr)



SIGN OF DERIVATIVE

○ Positive correlation

- More study time leads to higher score
- $r(s)$ score as a function of study time

$$\frac{dr}{ds} > 0 \quad \forall t, \quad r: \text{score}, s: \text{study time}$$

- One might get higher score by increasing the study time

○ Negative correlation

- Less party time leads to higher score
- $r(p)$ score as a function of party time


$$\frac{dr}{dp} < 0 \quad \forall p, \quad r: \text{score}, p: \text{party time}$$

- One might get higher score by decreasing the party time.



COMPUTATION: CHAIN RULE

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = g'(y)f'(x)$$




- Propagate the gradient along a path




EXAMPLE: LINEAR REGRESSION

- Solve it with gradient descent
- Linear regression: $\tilde{y} = w \cdot x + b$
- Loss: $L_{MSE} = E((y - \tilde{y})^2) \cong \sum_{i=1}^n (\tilde{y}_i - y_i)^2$



GRADIENT DESCENT FOR A NEURON

- $e(y) \equiv d - y$
- $y = g(h), h = w_1x_1 + w_2x_2$
- Objective: loss $L_{SE} = e^2$




$$\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial y} \frac{\partial y}{\partial h} \frac{\partial h}{\partial w_i} = (2e)(-1)g'(h)x_i$$

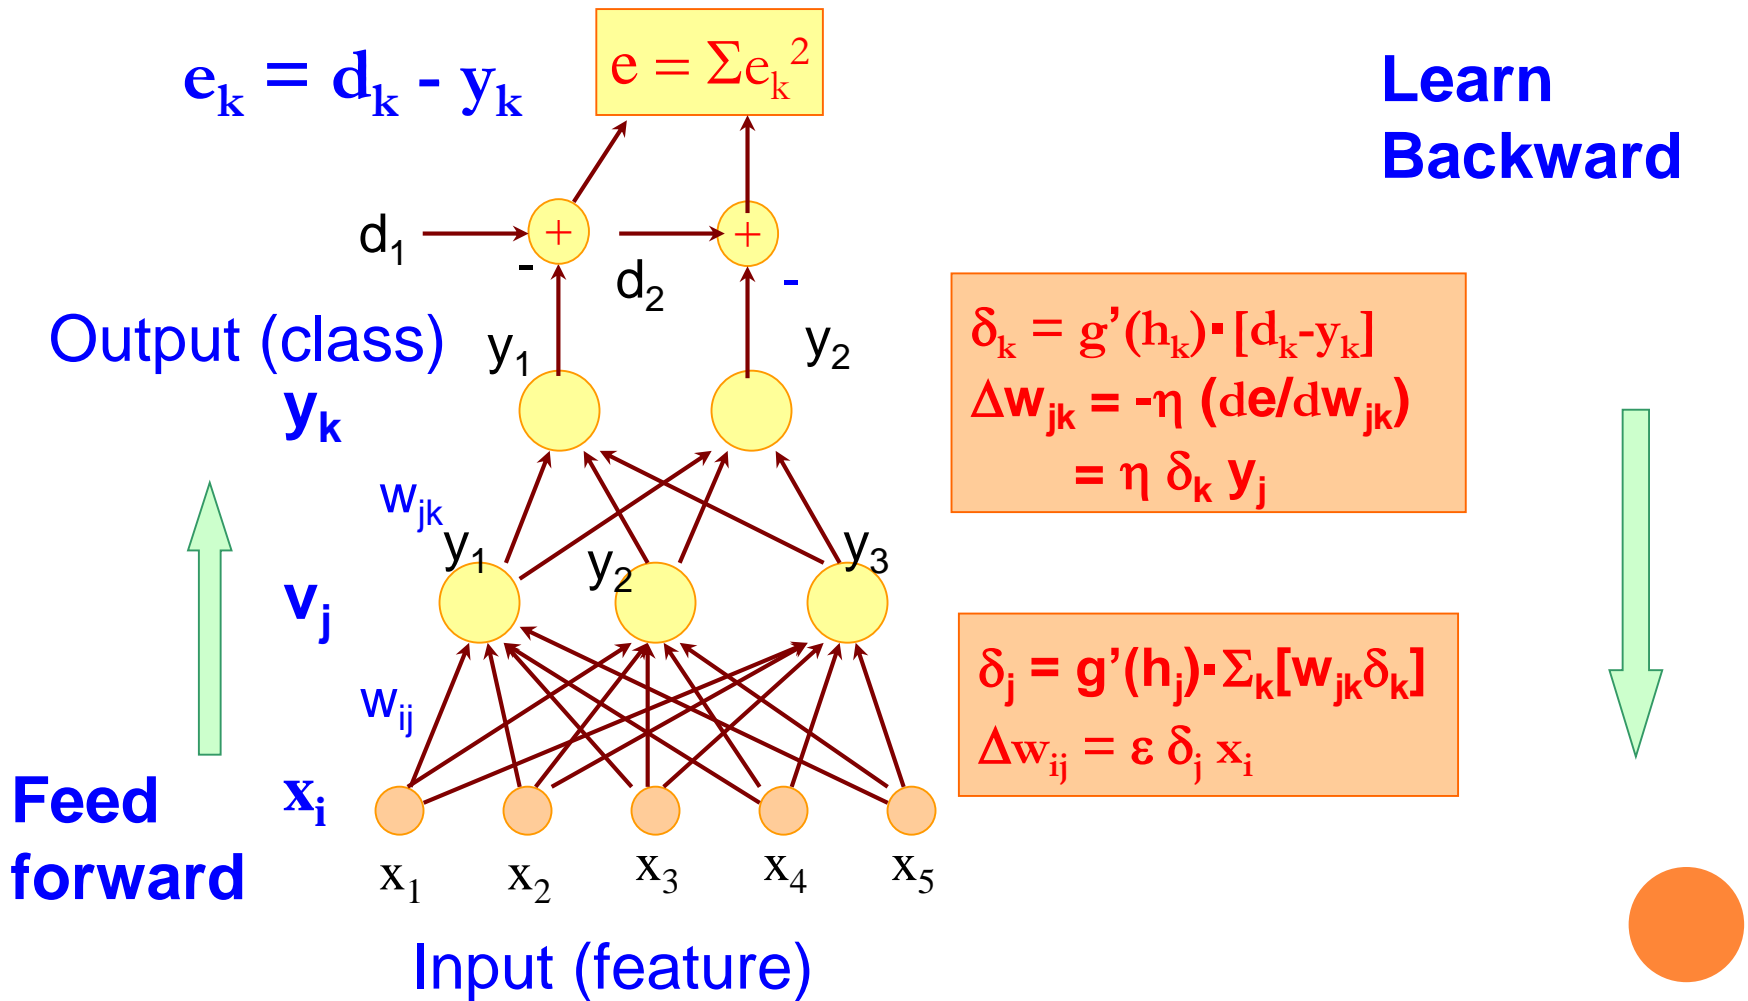
$$dw_i = -\varepsilon \frac{\partial L}{\partial w_i} = 2 \cdot \varepsilon \cdot g'(h) \cdot e \cdot x_i$$



SOLVING MLP



BACK PROPAGATION



BACK PROPAGATION

- i, j, k : neuron index for input layer (output x_i), hidden layer (output y_j), and output layer (output y_k) layers respectively.
- For node k in the output layer
 - the error e_k is the difference between the desired output and the actual output ($e_k = d_k - y_k$).
 - The error gradient for k is: $\delta_k = y_k \cdot (1 - y_k) \cdot e_k$
 - The weights are updated by: $w'_{jk} = w_{jk} + \alpha \cdot y_j \cdot \delta_k$
 - α : learning rate (a positive number below 1)
- Similarly, for a node j in the hidden layer
 - The error gradient for k is: $\delta_j = y_j \cdot (1 - y_j) \cdot \sum_k w_{jk} \delta_k$
 - The weights are updated by: $w'_{ij} = w_{ij} + \alpha \cdot x_i \cdot \delta_j$



LEARNING WEIGHTS OF OUTPUT LAYER

- $e = \sum_k e_k^2$, $e_k = d_k - y_k$

$$y_k = g(h_k), h_k = \sum_j w_{jk} y_j$$

by chain rule $\frac{de}{dw_{jk}} =$
 $(\frac{de}{de_k})(\frac{de_k}{dy_k})(\frac{dy_k}{dh_k})(\frac{dh_k}{dw_{jk}})$

$$= 2e_k (-1) \cdot g'(h_k) \cdot y_j$$


$$\Delta w_{jk} = -\alpha \left(\frac{de}{dw_{jk}} \right)$$

$$= -\alpha \cdot 2(d_k - y_k) (-1) \cdot g'(h_k) \cdot y_j$$

$$= \eta \cdot g'(h_k) \cdot (d_k - y_k) \cdot y_j$$

$$= \eta \cdot \delta_k \cdot y_j$$

$$\delta_k \equiv g'(h_k) \cdot (d_k - y_k) \text{ feedback of error}$$



LEARNING WEIGHTS OF HIDDEN LAYER

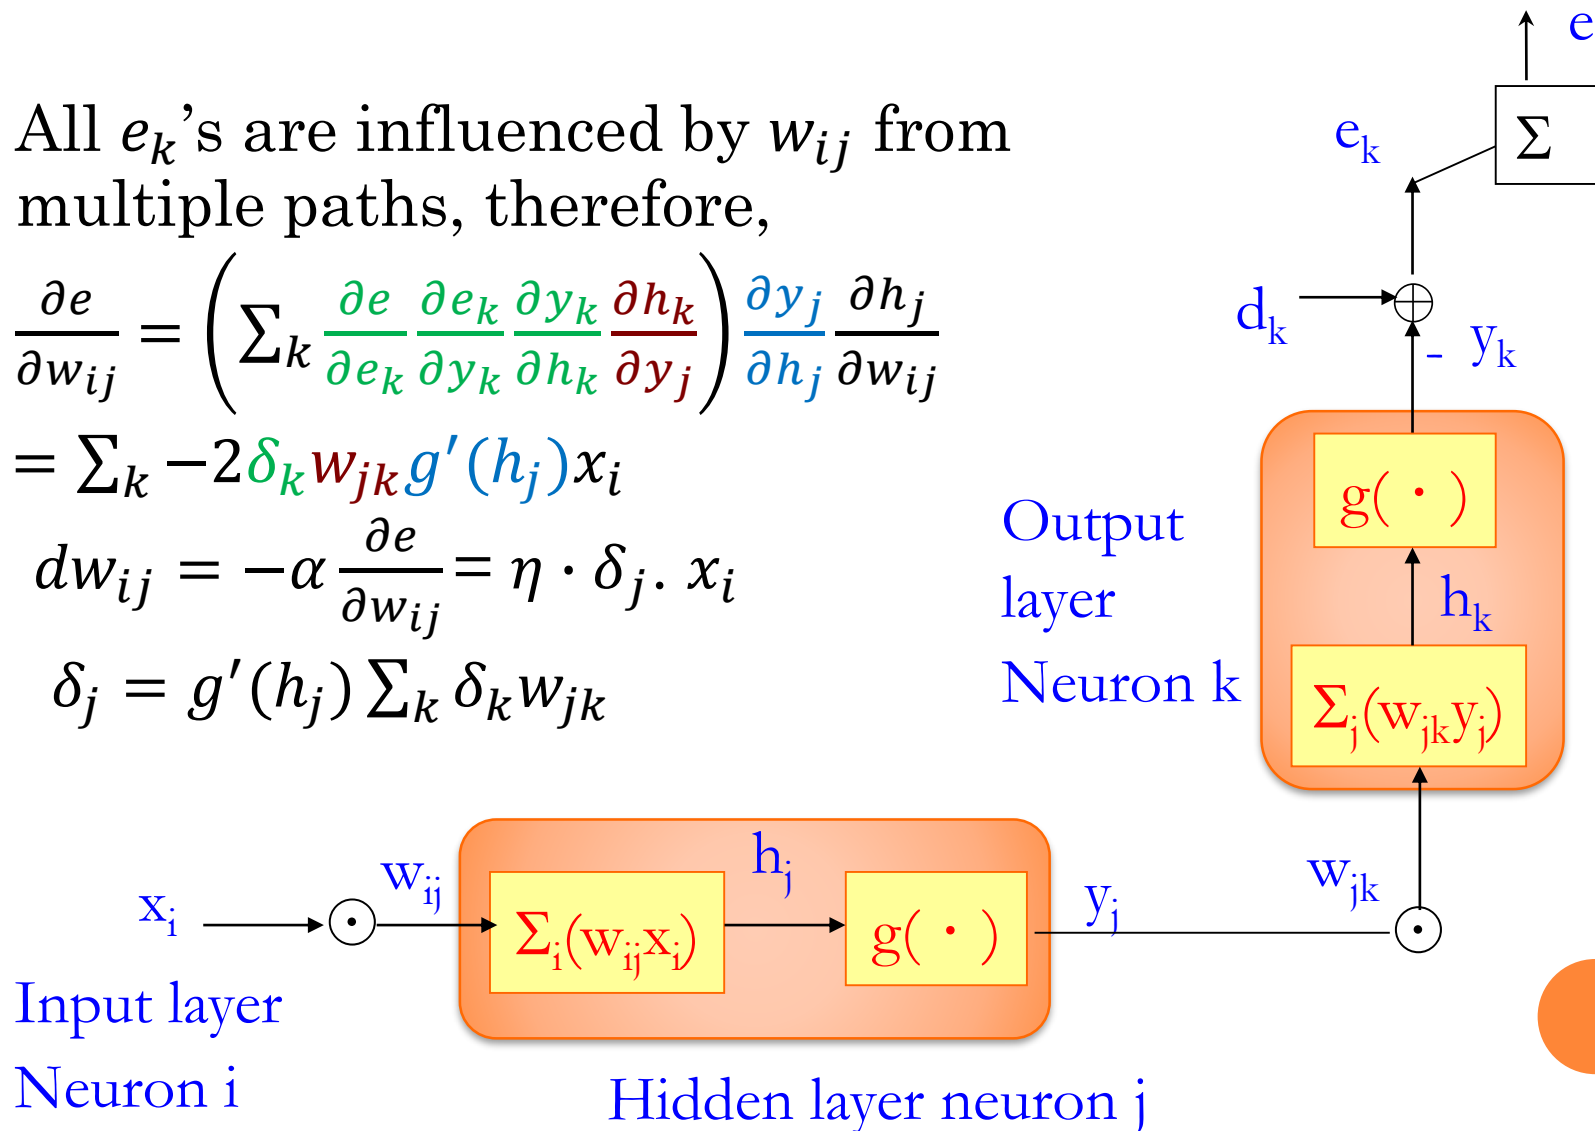
- All e_k 's are influenced by w_{ij} from multiple paths, therefore,

$$\frac{\partial e}{\partial w_{ij}} = \left(\sum_k \frac{\partial e}{\partial e_k} \frac{\partial e_k}{\partial y_k} \frac{\partial y_k}{\partial h_k} \frac{\partial h_k}{\partial y_j} \right) \frac{\partial y_j}{\partial h_j} \frac{\partial h_j}{\partial w_{ij}}$$


$$= \sum_k -2 \delta_k w_{jk} g'(h_j) x_i$$

$$\frac{\partial e}{\partial w_{ij}} = -\alpha \frac{\partial e}{\partial w_{ij}} = \eta \cdot \delta_j \cdot x_i$$

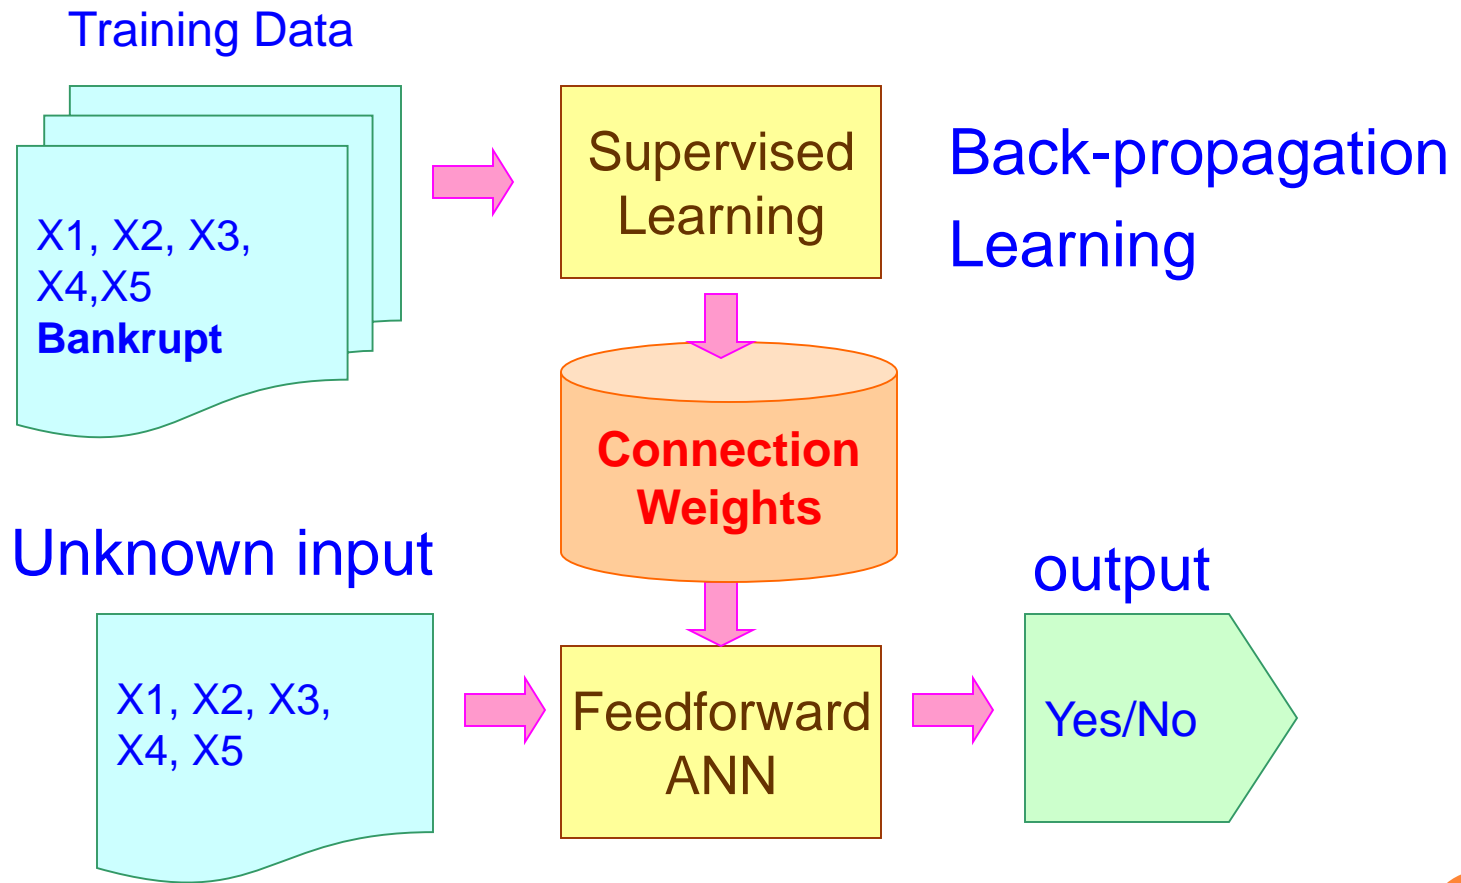
$$\delta_j = g'(h_j) \sum_k \delta_k w_{jk}$$



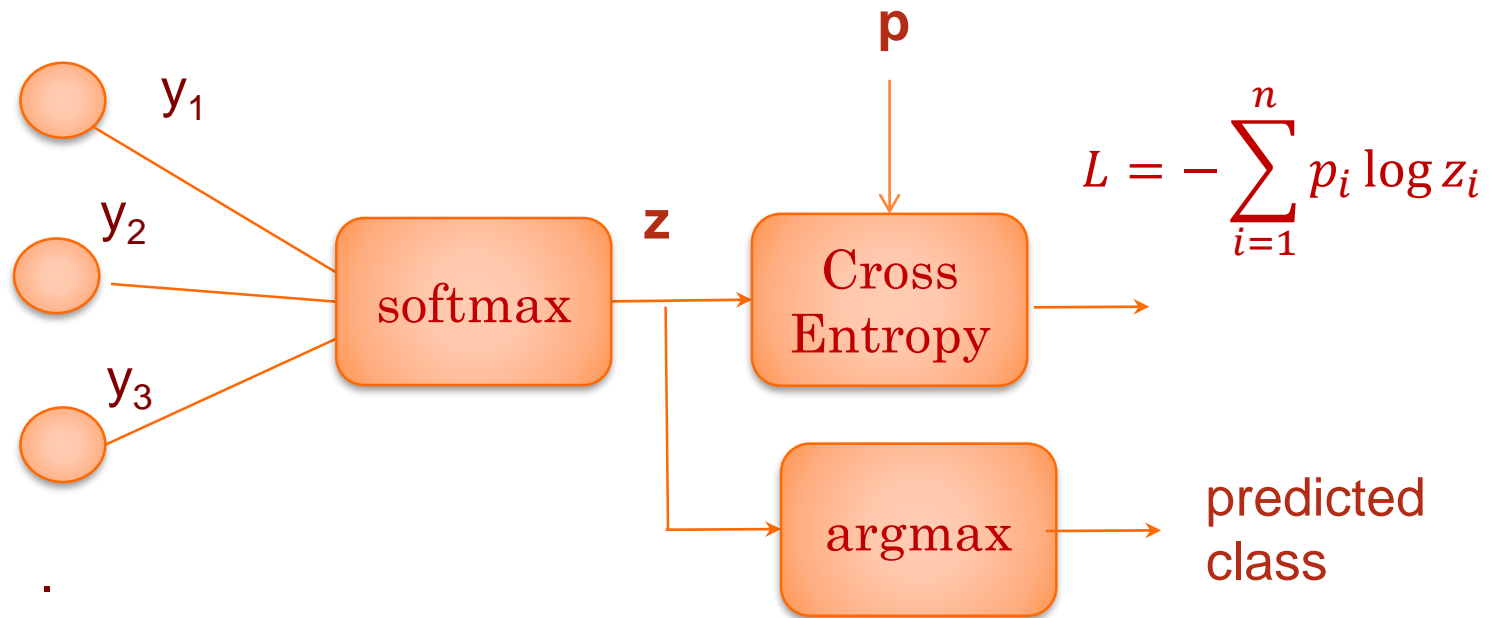
BINARY CLASSIFIER



MLP-BASED CLASSIFIER



SOFTMAX FOR M-ARY CLASSIFIER



- $z_i = \frac{e^{y_i}}{Z}$, $Z = \sum_{k=1}^n e^{y_k}$ where $\sum z_i = 1$
- p'_i s are the ground truths
- Cross entropy $L = - \sum_{i=1}^n p_i \log z_i$




GRADIENT PROPAGATION FOR SOFTMAX

- $z_i = \frac{e^{y_i}}{Z} = e^{y_i} Z^{-1}, Z = \sum_{k=1}^n e^{y_k}$
- $$\begin{aligned} \frac{\partial z_i}{\partial y_i} &= e^{y_i} Z^{-1} + e^{y_i} (-1) Z^{-2} \frac{\partial Z}{\partial y_i} \\ &= z_i - e^{y_i} Z^{-2} e^{y_i} = z_i - z_i^2 = z_i(1 - z_i) \end{aligned}$$

- $$\begin{aligned} \frac{\partial z_i}{\partial y_j} &= e^{y_i} (-1) Z^{-2} \frac{\partial Z}{\partial y_j} \\ &= -e^{y_i} Z^{-2} e^{y_j} = -z_i z_j \text{ for } j \neq i \end{aligned}$$

- $$G_{zy} = \begin{bmatrix} z_1(1 - z_1) & -z_1 z_2 & \dots & -z_1 z_n \\ -z_1 z_2 & z_2(1 - z_2) & \dots & -z_2 z_n \\ \vdots & \vdots & \ddots & \vdots \\ -z_1 z_n & -z_2 z_n & \dots & z_n(1 - z_n) \end{bmatrix}$$



MINIMIZING CROSS ENTROPY

- Objective: $L = -\sum_{i=1}^n p_i \log z_i$
- Let $\{z_i\}$ to approximate $\{p_i\}$ where $\sum p_i = 1$
- Since $\frac{\partial L}{\partial z_i} = -\frac{p_i}{z_i}$,

$$\nabla_{\mathbf{z}} L = \begin{bmatrix} \frac{\partial L}{\partial z_1} \\ \vdots \\ \frac{\partial L}{\partial z_n} \end{bmatrix} = - \begin{bmatrix} \frac{p_1}{z_1} \\ \vdots \\ \frac{p_n}{z_n} \end{bmatrix}.$$

$$\nabla_{\mathbf{y}} L = G_{\mathbf{zx}} \nabla_{\mathbf{z}} L = - \begin{bmatrix} z_1(1-z_1) & -z_1 z_2 & \dots & -z_1 z_n \\ -z_1 z_2 & z_2(1-z_2) & \dots & -z_2 z_n \\ \vdots & \vdots & \ddots & \vdots \\ -z_1 z_n & -z_2 z_n & \dots & z_n(1-z_n) \end{bmatrix} \begin{bmatrix} \frac{p_1}{z_1} \\ \vdots \\ \frac{p_n}{z_n} \end{bmatrix}$$



GRADIENT DESCENT ON SOFTMAX

- i -th element of \mathbf{g}_y is


$$\begin{aligned}\nabla_{y_i} &= -(z_i(1 - z_i) \frac{p_i}{z_i} - \sum_{j \neq i} z_i z_j \frac{p_j}{z_j}) \\ &= -(p_i(1 - z_i) - \sum_{j \neq i} z_i p_j) \\ &= -p_i + p_i z_i + z_i \sum_{j \neq i} p_j \quad (\sum p_j = 1) \\ &= -p_i + p_i z_i + z_i(1 - p_i) = \mathbf{z_i - p_i}.\end{aligned}$$

To minimize J , $-\nabla_{y_i} = (p_i - z_i)$ is propagated to y_i .

- Notice that classification is a special case
 - $p_k = 1$ and $p_j = 0$ for all $j \neq k$ for gold branch k (one-hot)
 - i.e. $1 - z_k$ for target node and $-z_j$ for the others.
- \mathbf{z} is an estimate of the target distribution, \mathbf{p} .



SOFTMAX: PROPAGATING GRADIENTS



IMPROVING SPEED OF CONVERGENCE

- Use **generalized delta rule** for learning
 - $\Delta w_{ij}(t) = \alpha \cdot x_i \cdot \delta_j + \beta \cdot \Delta w_{ij}(t-1)$
 - **Momentum** β has a typical value of 0.95
 - Weighted sum between training data and original value (learn incrementally)
- Use **hyperbolic tangent function** instead of sigmoid
 - $\tanh(x) = 2a/(1+e^{-bx}) - a$ ($a = 1.7$, $b = 0.7$)
- Use **dynamic learning rate** α
 - Increasing α if square-error changes in the same direction for several epochs, and decreasing α otherwise.



ADAPTIVE LEARNING RATE

○ AdaGrad

- $\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,i} + \epsilon}} g_{t,i}$
- $g_{t,i} = \nabla_{\theta} J(\theta_{t,i})$
- ϵ is smoothing value (typically 1e-8)
- $G_{t,i}$ is energy of gradients: summation of the squares of past gradients ($g_{t,i}^2$)

○ AdaDelta

- $E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma) g_t^2$
- $\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$
- Incremental estimation of average energy $E[g^2]_t$
- Momentum $\gamma=0.99$ typically



ADAM (ADAPTIVE MOMENT ESTIMATION)

- $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
- $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$
- $\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{v_t + \epsilon}} m_t$
 - It is required to normalize m_t and v_t by
$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \hat{v}_t = \frac{v_t}{1 - \beta_2^t} \text{ if } m_0 = v_0 = 0$$
 - But unnecessary if $m_0 = g_0, v_0 = g_0^2$
- Exponential decay ($\beta < 1$) of influence
 - $m_1 = \beta_1 g_0 + (1 - \beta_1) g_1$
 - $m_2 = \beta_1(\beta_1 g_0 + (1 - \beta_1) g_1) + (1 - \beta_1) g_2$
$$= \beta_1^2 g_0 + \beta_1 (1 - \beta_1) g_1 + (1 - \beta_1) g_2$$
 - $m_t = \beta_1^t g_0 + \beta_1^{t-1} (1 - \beta_1) g_1 + \dots + (1 - \beta_1) g_t = \sum_{i=0}^t w_i g_i$
 - $\sum_{i=0}^t w_i = 1 \rightarrow m_t$ as weighting sum of gradients



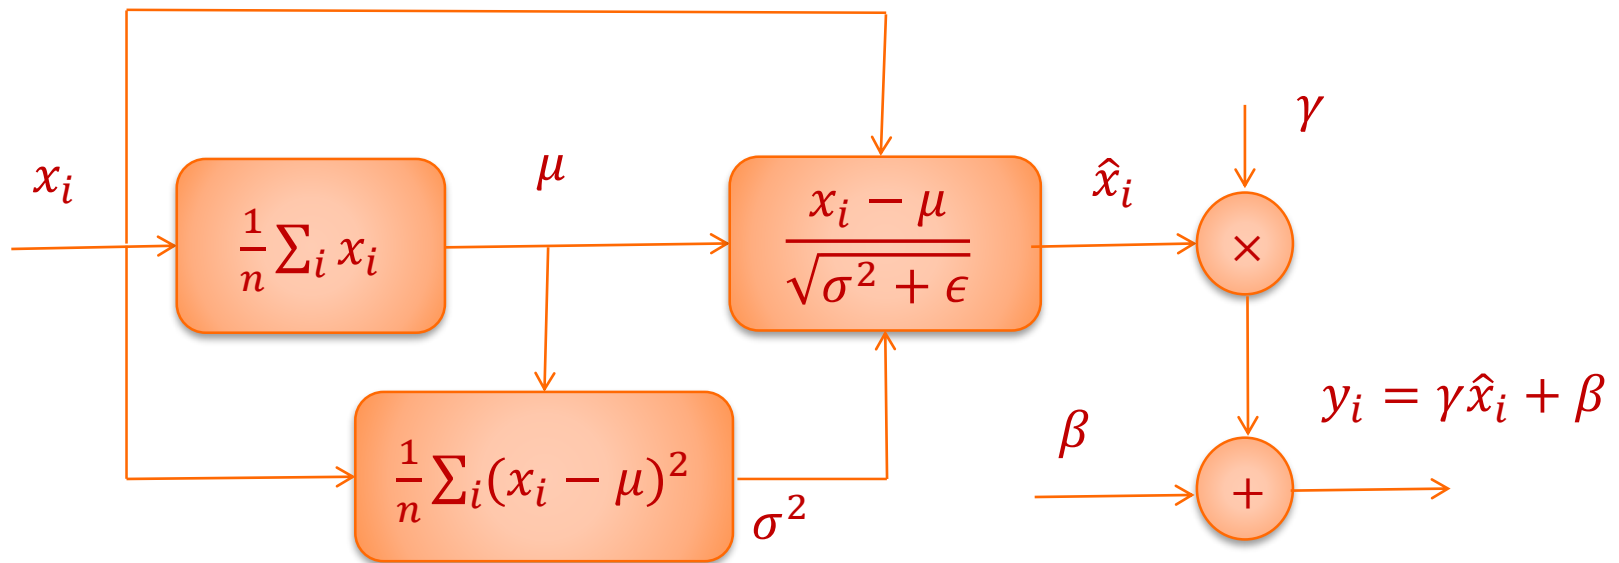
BATCH NORMALIZATION

- Adjust the distribution (means and variances) of neuron inputs to avoid gradient vanishing (especially for deep network)
- May accelerate the convergence when training the network
- Reference:

“Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”, Sergey Ioffe and Christian Szegedy, 2015.




BATCH NORMALIZATION



- \hat{x}_i has zero mean and unit variance.
- y_i with shifted mean and scaled variance.
- γ and β are learnable.
- When β equals to mean and γ equals to standard deviation, BN is identity transform.



BATCH NORMALIZATION: GRADIENT



NORMALIZATION

- $z_i = \frac{y_i}{Y} = y_i Y^{-1}, Y = |\mathbf{y}| = S^{1/2}$

$$S = \sum_{k=1}^n y_k^2$$


- $$\begin{aligned} \frac{\partial z_i}{\partial y_i} &= Y^{-1} + y_i(-1)Y^{-2} \frac{\partial Y}{\partial y_i} \\ &= Y^{-1} - y_i Y^{-2} \left(\frac{1}{2}\right) S^{-\frac{1}{2}} (2y_i) \\ &= Y^{-1} (1 - y_i Y^{-2} y_i) = Y^{-1} (1 - z_i^2) \end{aligned}$$

- $$\begin{aligned} \frac{\partial z_i}{\partial y_j} &= y_i(-1)Y^{-2} \frac{\partial Y}{\partial y_j} = -y_i Y^{-2} \left(\frac{1}{2}\right) S^{-\frac{1}{2}} (2y_j) \\ &= -y_i Y^{-2} Y^{-1} y_j = -Y^{-1} z_i z_j \end{aligned}$$

- Define G_{zy} such that

$$G(i, i) = 1 - z_i^2$$

$$G(i, j) = -z_i z_j$$



RECURRENT NETWORKS


- Feed forward networks do not have memory.
- Recurrent networks can have connections between nodes in any layer, which enables them to store data – a **memory**.
 - flip-flop
- Capability of sequence prediction
 - Recurrent networks can be used to solve problems where the solution depends on previous inputs as well as current inputs (e.g. predicting stock market movements).



HOPFIELD NETWORK

- $\underline{\xi} = (\xi_1, \xi_2, \dots)$ to be memorized with values of +1/-1
- $\xi_i = \text{sign}(\sum_j w_{ij} \xi_j - \theta_i)$
 $w_{ij} = (1/N) \xi_i \xi_j$
- $h_i = \sum_j w_{ij} \xi_j = (1/N) \sum_j \xi_i \xi_j \xi_j$
 $= \xi_i$ (since $\xi_j^2 = 1$)
- $h_i = \sum_j w_{ij} S_j = (1/N) \sum_j \xi_i \xi_j S_j$
 $= (1/N) \xi_i (\sum_j \xi_j S_j)$
 $= (1/N) \xi_i \cdot \langle \underline{S}, \underline{\xi} \rangle$

If \underline{S} is close to $\underline{\xi}$ (in Hamming distance) such that $\langle \underline{S}, \underline{\xi} \rangle > 0$, then S_i will be attracted into ξ_i and ($\underline{S} = \underline{\xi}$)



HOPFIELD NETWORK (CONT'D)

$$W = \sum_{i=1}^N X_i X_i^t - N I$$

$$Y = \text{sign}(WX - \theta)$$


X_i 's are patterns to be memorized

Weights: N patterns (X_i 's)

$$Y_i = \text{Sign}(WX_i - \theta) = X_i$$

Input X close to X_i will be attracted to X_i

Hopfield network is a memory that usually maps an input vector to the memorized vector whose **Hamming distance** from the input vector is least. (auto-associative memory) (memory \rightarrow recall)



EXAMPLE OF HOPFIELD NETWORK

$$X_1 = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} \quad X_2 = \begin{bmatrix} -1 \\ -1 \\ -1 \\ -1 \\ -1 \end{bmatrix} \quad X_3 = \begin{bmatrix} 1 \\ -1 \\ 1 \\ 1 \\ -1 \end{bmatrix}$$

$$W = \sum_{i=1}^3 X_i X_i^t - 3I = \begin{bmatrix} 0 & 1 & 3 & 3 & 1 \\ 1 & 0 & 1 & 1 & 3 \\ 3 & 1 & 0 & 3 & 1 \\ 3 & 1 & 3 & 0 & 1 \\ 1 & 3 & 1 & 1 & 0 \end{bmatrix}$$

$$Y_1 = \text{sign}(WX_1) = \text{sign}\left(\begin{bmatrix} 8 \\ 6 \\ 8 \\ 8 \\ 6 \end{bmatrix}\right) = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = X_1$$

$$Y_2 = WX_2 = X_2 \quad Y_3 = WX_3 = X_3$$

$$X_4 = \begin{bmatrix} 1 \\ 1 \\ -1 \\ 1 \\ 1 \end{bmatrix}$$

$$Y_4 = \text{sign}(WX_4) = \text{sign}\left(\begin{bmatrix} 2 \\ 4 \\ 8 \\ 2 \\ 4 \end{bmatrix}\right) = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = X_1$$




EXAMPLE OF HOPFIELD NETWORK (CONT'D)

$$X_5 = \begin{bmatrix} -1 \\ 1 \\ -1 \\ 1 \\ 1 \end{bmatrix}$$

$$Y_5 = \text{sign}(WX_5) = \text{sign}\left(\begin{bmatrix} 2 \\ 2 \\ 2 \\ -4 \\ 2 \end{bmatrix}\right) = \begin{bmatrix} 1 \\ 1 \\ 1 \\ -1 \\ 1 \end{bmatrix}$$

$$Y_5' = WY_5 = \text{sign}\left(\begin{bmatrix} 2 \\ 4 \\ 2 \\ 8 \\ 4 \end{bmatrix}\right) = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = X_1$$



Auto-associative Memory

Storage: setting the weights

Retrieval: retrieving data



APPLICATION OF HOPFIELD NETWORK

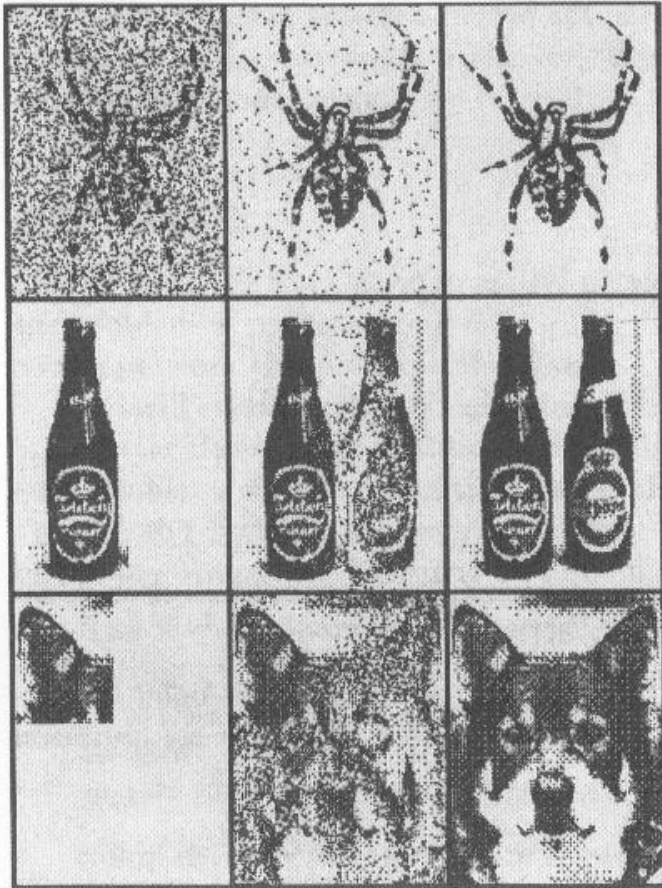


FIGURE 2.1 Example of how an associative memory can reconstruct images. These are binary images with 130×180 pixels. The images on the right were recalled by the memory after presentation of the corrupted images shown on the left. The middle column shows some intermediate states. A sparsely connected Hopfield network with seven stored images was used.




BIDIRECTIONAL ASSOCIATIVE MEMORIES (BAM)

$$W = \sum_{i=1}^n X_i Y_i^t$$

$$\text{sign}(W^t X_j) = Y_j$$

$$\text{sign}(W Y_j) = X_j$$

- association $X \rightarrow Y$



EXAMPLE OF BAM

$$X_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \quad X_2 = \begin{bmatrix} -1 \\ -1 \end{bmatrix} \quad Y_1 = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \quad Y_2 = \begin{bmatrix} -1 \\ -1 \\ -1 \end{bmatrix}$$

$$W = \sum_{i=1}^2 X_i Y_i^t = \begin{bmatrix} 2 & 2 & 2 \\ 2 & 2 & 2 \end{bmatrix}$$


$$\text{sign}(W^t X_1) = \text{sign}\left(\begin{bmatrix} 4 \\ 4 \\ 4 \end{bmatrix}\right) = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = Y_1$$

$$\text{sign}(W Y_1) = \text{sign}\left(\begin{bmatrix} 6 \\ 6 \end{bmatrix}\right) = \begin{bmatrix} 1 \\ 1 \end{bmatrix} = X_1$$



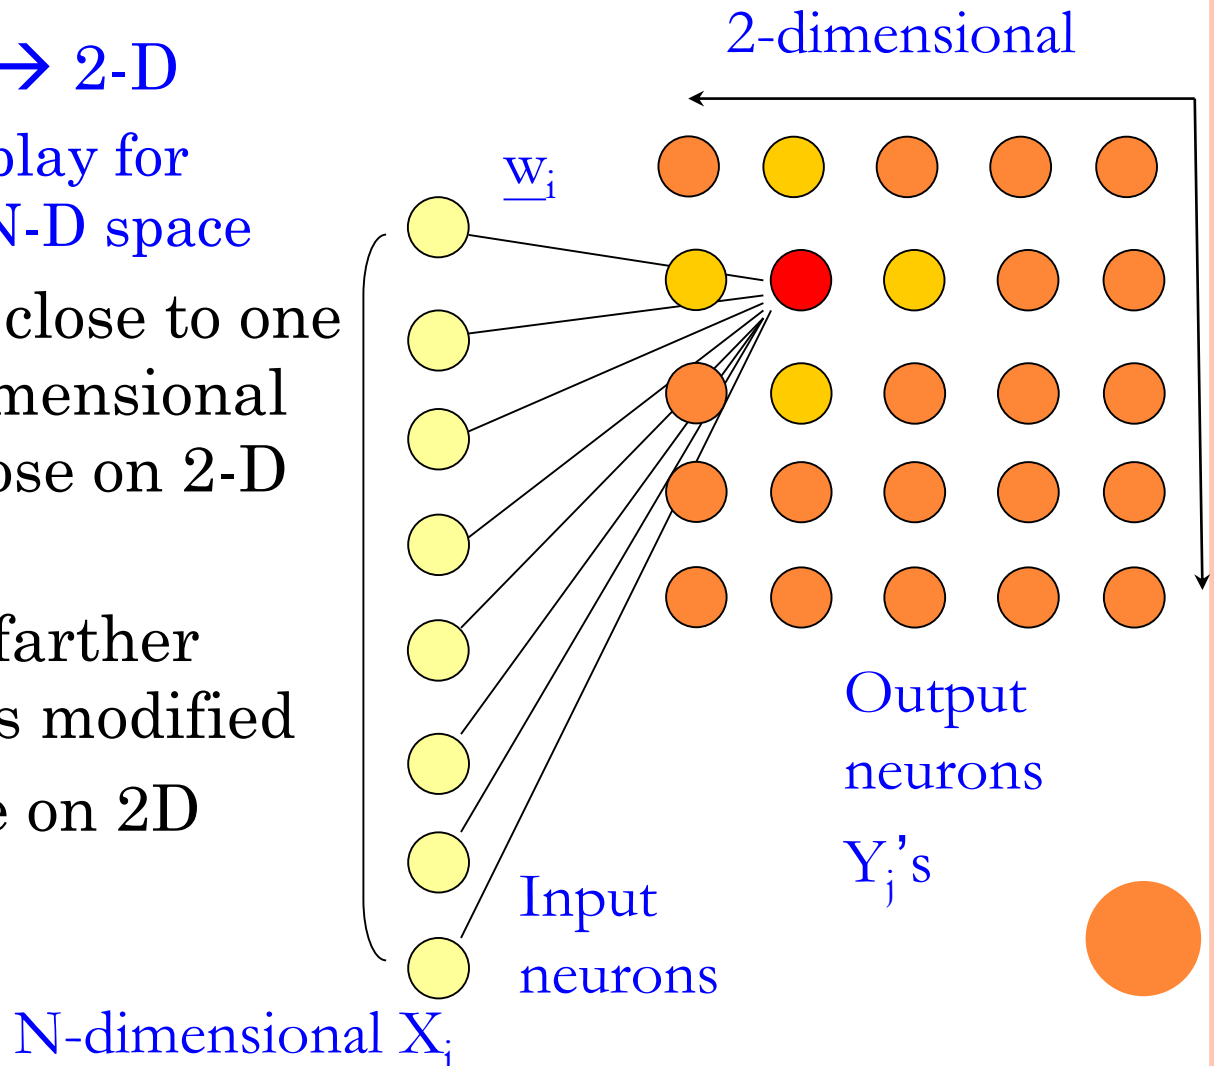
KOHONEN MAPS (SELF-ORGANIZING MAP)

- Unsupervised learning
 - Clustering for training data $\{\underline{x}_j\}$
- Winner-take-all
 - The winner is the neuron with highest output (\underline{w}_i closest to \underline{x})
 - The **winner and its spatial neighborhoods** can update weights
 - $w_{ij}' = w_{ij} + \alpha(x_j - w_{ij})$ or $d\underline{w}_i = \alpha(\underline{x} - \underline{w}_i)$
- $y_i = \text{sign}(\langle \underline{x}, \underline{w}_i \rangle)$
 - $\langle \underline{x}, \underline{w}_i \rangle$ similarity
- 和KNN相似, K個output neurons的 weights代表K個centroids



KOHONEN MAPS (CONT'D)

- N-dimensional \rightarrow 2-D
 - 2D Visual display for closeness for N-D space
- Codewords (\underline{w}_i) close to one another in n-dimensional space will be close on 2-D plane
- The weights of farther neurons are less modified
- View N-D space on 2D plane



REFERENCES

- *Introduction to the Theory of Neural Computing*
John Hertz, Anders Krogh, Richard G. Palmer
- *Artificial Intelligence A modern Approach*
Stuart Russel, Peter Norvig
- *Artificial Intelligence, illuminated*
Ben Coppin
- *An overview of gradient descent optimization algorithms*, Sebastian Ruder, 2016.

