



HIDDEN MARKOV MODEL

Bor-shen Lin

bslin@cs.ntust.edu.tw

<http://www.cs.ntust.edu.tw/~bslin>

OBSERVATIONS FROM RANDOM PROCESS

- **Observable output** generated by real-world random process
- Discrete vs. Continuous
 - Character, gender, class, ...
 - Speech, temperature,...

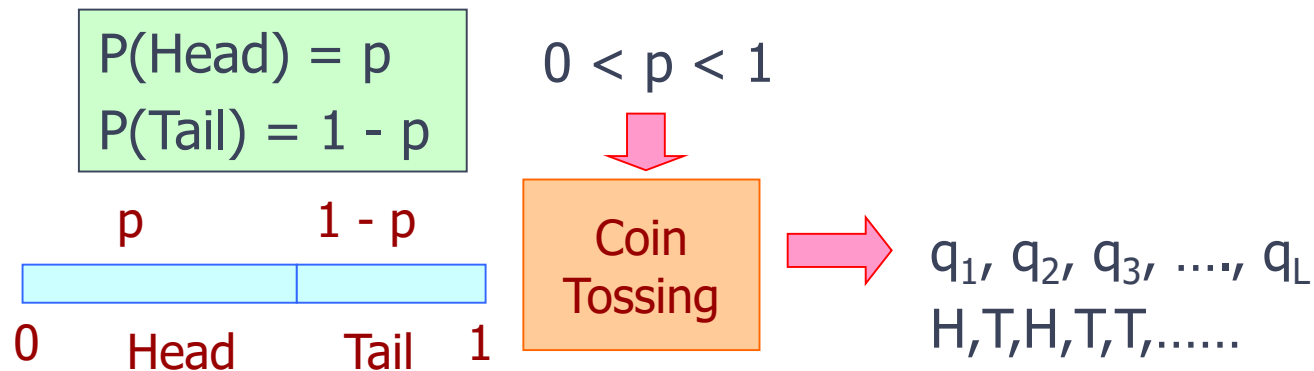


HIDDEN MARKOV MODEL

- A type of *stochastic model*
- First published in mathematic journals
 - Baum and his colleagues, late 1960s
- Applied to speech recognition
 - Baker at CMU/Jelinek at IBM, 1970s
- A general approach
 - Can be applied to many other areas



A SIMPLE STOCHASTIC PROCESS



- The outcomes q_1, q_2, \dots, q_T can be obtained through random tests
 - q_i has Two states: Head (H) and Tail (T)
- Each tossing is *independent* of the results of previous tossings
 - Random variables q_1, q_2, \dots, q_L are i.i.d.
- **States** (H,T): possible **outcomes** of random variables q_t

DICE TOSSING

- Fair dice: $(1/6, 1/6, 1/6, 1/6, 1/6, 1/6)$ for all faces
- Every tossing is **independent** of the others
 - Current tossing does not depend on previous results
- Number of states : 6
 - $q_t = S_i, i = 1, 2, \dots, 6$
 - q_t random variables
 - S_i outcome symbols (number of points)




DICE TOSSING (CONT'D)

- 6 dices numbered from 1 to 6
 - With different distributions (probably unfair)
 - Dice 1: (1/3, 1/10, ...)
 - Dice 2: (1/4., 1/6, ...)
 - ...
 - Dice 6: (1/5, 1/8, ...)
 - The choice of dice **depends on** the outcome of previous tossing
 - E.g. if $q_t = S_4$ (4 points) in t_{th} tossing, then choose Dice 4 for $(t+1)_{th}$ tossing
 - q_t **depends on** q_{t-1}
 - $q_1 \rightarrow q_2 \rightarrow q_3 \rightarrow \dots \rightarrow q_T$



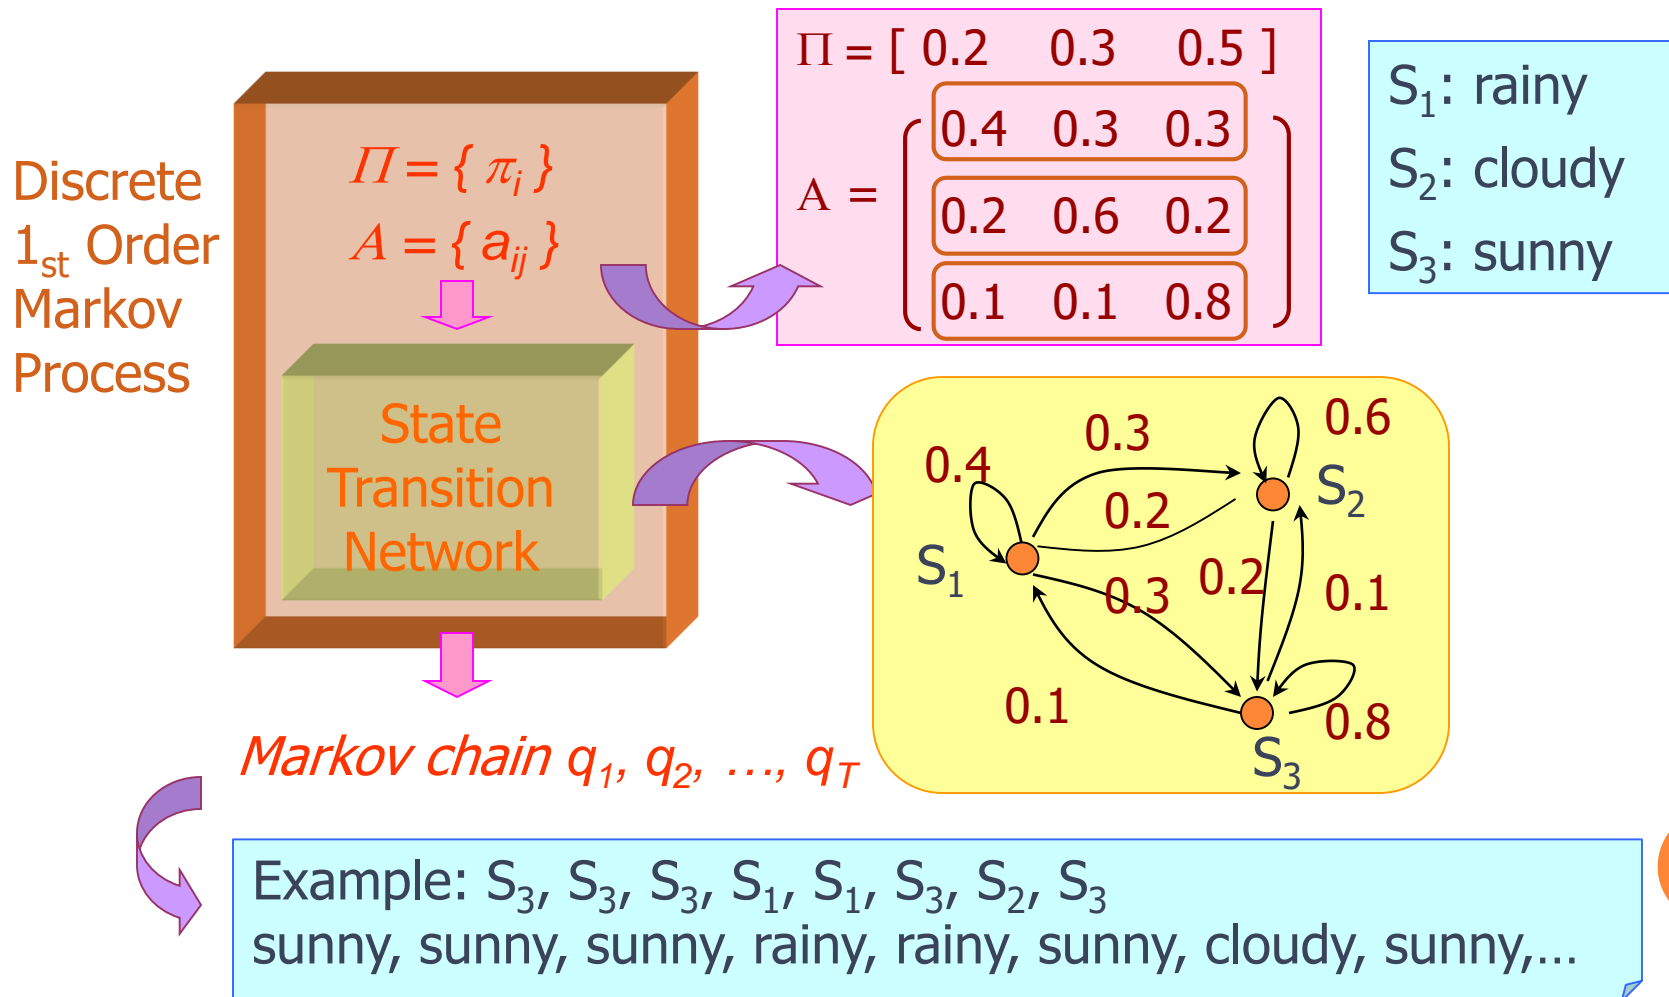
DISCRETE MARKOV PROCESS



A system with 5 states $S_1 - S_5$.

- System is in one of N distinct states (S_1, S_2, \dots, S_N) at any time, and *state sequence* q_1, q_2, \dots, q_T is observed
- First-order Markov Chain
 - $P(q_t=S_j|q_{t-1}=S_i, q_{t-2}=S_k \dots) = P(q_t=S_j|q_{t-1}=S_i)$
 - $a_{ij} \equiv P(q_t=S_j|q_{t-1}=S_i) \quad 1 \leq i, j \leq N$
 $a_{ij} > 0, \sum_{j=1}^N a_{ij} = 1$
 - $\pi_i \equiv P(q_1=S_i) \quad 1 \leq i \leq N$
- State transition depends on *only the previous state*
- State transits *stochastically*
- Imagine: N dices with N faces each


MARKOV PROCESS SIMULATOR




MODEL VS. REALITY

○ Model \neq Reality

- Does God use Markov Process to generate weather patterns?
- Real-world is usually *unknown* and thus *approximated* by some model
- Amount of parameters may influence the accuracy



Area of a lake



DISCUSSIONS

○ Simulation

- It is possible to simulate a (true) Markov process, and generate random observations accordingly

○ Modeling the real world

- In modeling the true random process in real world, we have no idea whether the true (unknown) random process is Markov process or not.

We simply model it!

- The observations can be used to estimate the model parameters such that the model can best account for the statistics of the observations



Estimation of Model Parameters

③ $\hat{a}_{ij} = \frac{c_{ij}}{\sum_k c_{ik}}$

c_{ik} is the occurrence count of transitions from state i to state k in all observation sequences

○ $\hat{\pi}_j = \frac{c_j}{\sum_k c_k}$


c_k is the occurrence count that the first state in a state sequence is state k for all observation sequences

○ Example

- $S_1, S_2, S_2, S_3, S_1, S_1, S_3, \dots$
- $c_{12}^{++}, c_{22}^{++}, c_{23}^{++}, c_{31}^{++}, c_{11}^{++}, c_{13}^{++}, \dots$



STATE DIAGRAM



State sequence $Q = \{ q_t \}$

$q_1, q_2, q_3, q_4, q_5, \mathbf{q}_6, q_7, q_8, \dots, q_T$

$S_3, S_3, S_3, S_1, S_1, \mathbf{S}_3, S_2, S_3, \dots$




PROBABILITY OF STATE SEQUENCE

- The probability of Markov model (Π, A) generating the state sequence $Q = S_3, S_3, S_3, S_1, S_1, S_3, S_2, S_3$

- $P(Q|\Pi, A) = P(S_3) \cdot P(S_3|S_3) \cdot P(S_3|S_3) \cdot P(S_1|S_3) \cdot$

$$P(S_1|S_1) \cdot P(S_3|S_1) \cdot P(S_2|S_3) \cdot P(S_3|S_2)$$

$$= \pi_3 \cdot a_{33} \cdot a_{33} \cdot a_{31} \cdot a_{11} \cdot a_{13} \cdot a_{32} \cdot a_{23}$$



- System stays in S_i for time D_i (random variable)

- $P(D_i = d|\Pi, A) = P(Q = \overbrace{S_i, S_i, S_i, \dots, S_i}^d, S_{j \neq i} | \Pi, A, q_1 = S_i)$

$$= (a_{ii})^{d-1} (1 - a_{ii}) \equiv p_i(d) \quad d = 1, 2, \dots$$

- Expectation $E(D_i) = \sum_d (d \cdot p_i(d)) = 1 / (1 - a_{ii})$ $a_{ii} \uparrow, E(D_i) \uparrow$

- Transition matrixes: modeling the duration

FROM MARKOV MODEL TO HMM

○ Markov Model


- State transits stochastically and state sequence can be observed

○ *Hidden* Markov Model

- State transits stochastically, but state sequence *cannot be observed*
- State sequence is unseen $\{ q_t \}$
 - rainy, cloudy, sunny
- Observation sequence $\{ O_t \}$
 - very hot (VH), hot(H), warm(W), cool(CL), cold(CD)



URN AND BALL MODEL



- State transition
 - Roll the dice to select a urn
 - Multiple dices
- Observation
 - Pick up a ball from the selected urn
- The urns(j) has an associated distribution of balls, $b_j(O_t)$
- Double stochastic processes



ELEMENTS OF HMM

- N : number of states
 - S_1, S_2, \dots, S_N (outcome symbols, not random)
- Initial state distribution $\Pi = \{\pi_i\}$
 - $\pi_i = P(q_1 = S_i)$
- State transition distribution $A = \{a_{ij}\}$
 - $a_{ij} = P(q_{t+1} = S_j | q_t = S_i)$
- State observation distribution $B = \{b_j(O_t)\}$
 - $b_j(O_t) = P(O_t | q_t = S_j)$
 - The output O_t depends on the current state q_t
 - Could be *discrete* or *continuous* functions
- $\lambda = (\Pi, A, B)$



OBSERVATION DISTRIBUTION

- Discrete : Probability weighting function
 - $b_j(O_t = \nu_k)$ $k = 1, 2, \dots, K$
 - ν_k may be obtained through **encoding/quantization**
 - Ex. $O_t = \nu_k$: very hot, hot, warm, cool, cold, very cold
- Continuous : Probability density function
 - $b_j(O_t = x)$ could be a continuous function in any region which satisfied $\int b_j(O_t = x) dx = 1$
 - Ex. temperature could be 34, 31.5, 28, 29.3,
 - Example: weighted Gaussian mixtures
 - $b_j(O_t = x) = \sum_{m=1} c_{jm} \cdot \mathcal{N}(x; m_{jm}, \sigma_{jm}^2)$ $-\infty < x < \infty$

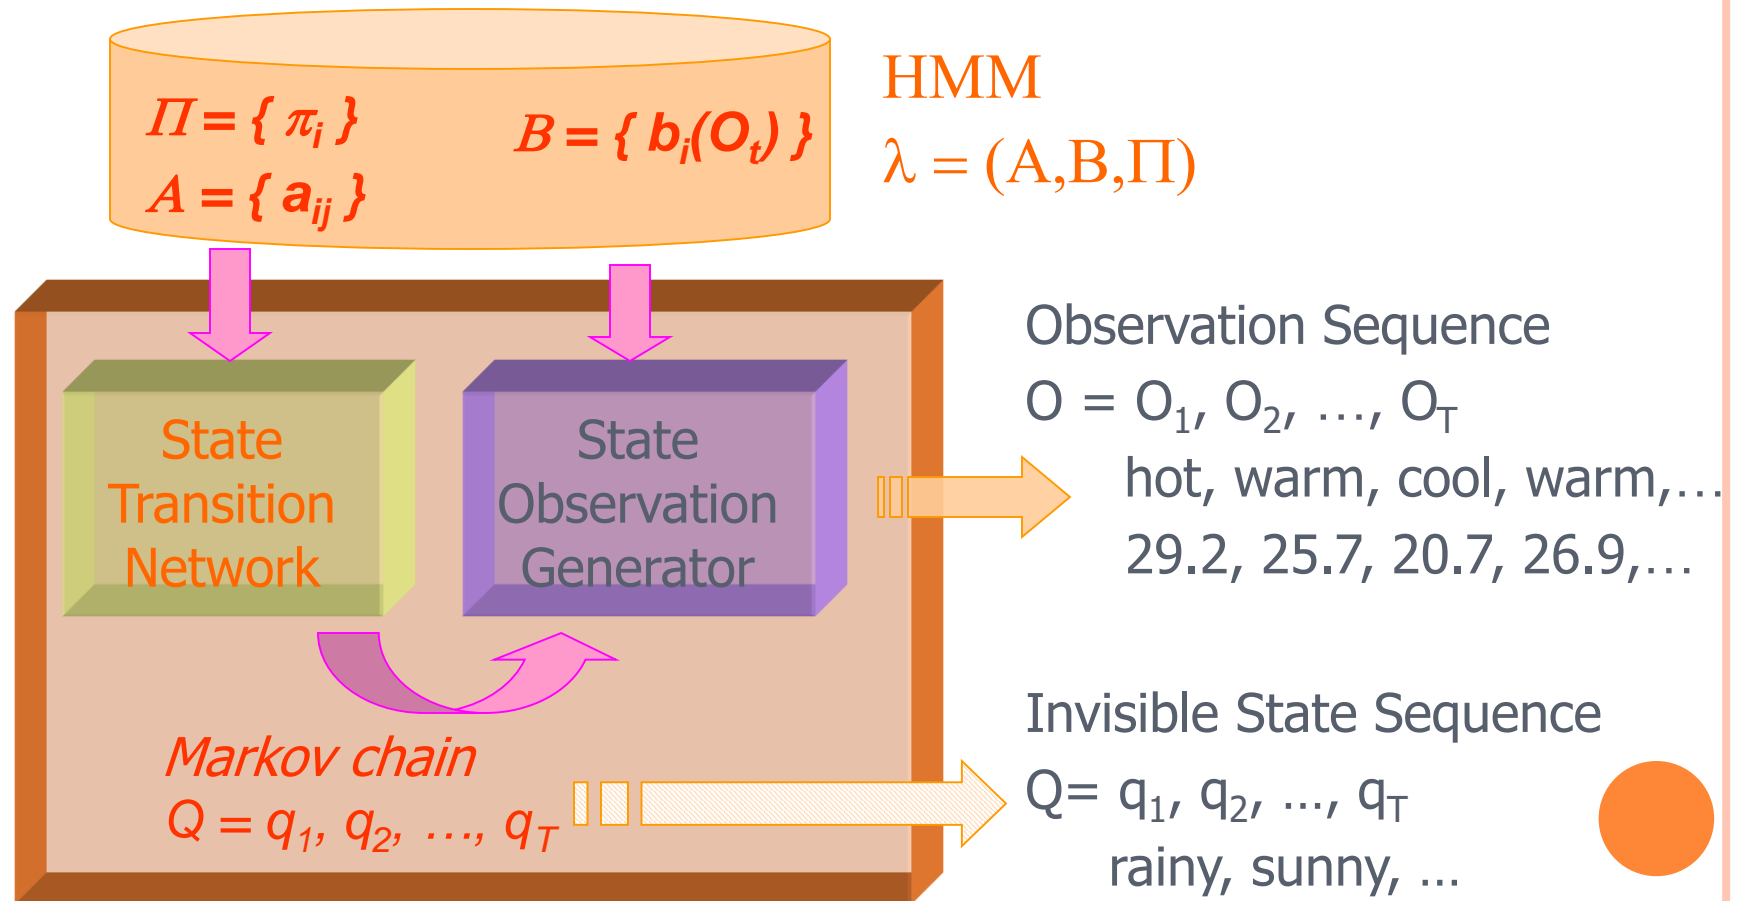


SEMI-CONTINUOUS OBSERVATION DISTR.

- Discrete : $b_j(O_t = v_k) = b_{jk}$
 - Quantization error, insufficient training (e.g. $b_{kj} = 0$)
- Continuous: $\mathcal{N}(m_{jm}, \sigma_{jm}^2)$ for state j & mixture m
 - Computation load is dependent of the numbers of models, states, or mixtures
 - 500 HMMs, 4 states, 10 mix \rightarrow 20000 computations!
- Semi-Continuous : $b_j(O_t=x) = \sum_{k=1} b_{jk} \cdot \mathcal{N}(x; m_k, \sigma_k^2)$
 - Common Gaussian mixtures are shared among all states, but with all states have different mixture weights
 - Trade-off between discrete and continuous distribution
 - Can reduce the effect of **quantization errors** and insufficient training due to parameter sharing
 - m_k and σ_k are obtained from observations clustered to codeword v_k (codebook size is fixed)



HMM SIMULATOR



CONCEPTS OF HIDDEN STATES

- Economic conditions
 - Hidden States: economic conditions
 - Observation: economic indexes
- Moods
 - Hidden States: mood states
 - Observation: color of dressing, face expression, interaction style
- Periods of life
 - Hidden States: stages of growth
 - Observation: weight, height, and other features
- Seasons
 - 4 seasons



TO EVALUATE THE LEARNED

- Given HMM λ and unknown observation O
 - Calculation of $P(O | \lambda)$
 - $i^* = \operatorname{argmax}_i P(O | \lambda_i)$: to find the HMM that most probably generates O
- Use for Recognition
 - With observed sequence O and the HMMs for some λ_i 's, calculate $P(O | \lambda_i)$
 - Decide which model is most probable to generate the unknown observation O



TO LEARN FROM OBSERVATION

- We have the observation O , and *assume* it is generated by a HMM.
 - Is it possible to *learn/train* the *HMM parameters* λ that is optimal in some sense from the observation O ?
 - HMM λ^* can *best represent* the observation O
 - The stochastic characteristics of observation O are caught and kept in model parameters of λ^*
 - $O \rightarrow \lambda^* = (\Pi, A, B)$
- Example
 - Temperature sequence $O \rightarrow \text{HMM } \lambda^*$



TO GUESS THE UNSEEN STATES Q

- Assume observation O is generated by an HMM λ .
 - Is it possible to reconstruct the *state sequence* Q^* that is optimal in some sense?
- Example
 - Temperature (O) & HMM (λ) \rightarrow weather (Q^*)
 - C, H, H, VH, ... \rightarrow rainy, sunny, sunny, sunny, ...




FORMULATE PROBLEMS

- Evaluating: $P(O | \lambda)$
- Learning: $O \rightarrow \lambda^*$
- Labeling: $O, \lambda \rightarrow Q^*$
- $O = O_1, O_2, \dots, O_T$ (random variables)
- $Q = q_1, q_2, \dots, q_T$ (random variables)
- λ : model parameters
- Important Algorithms
 - Forward-backward procedure
 - Expectation-Maximization (EM) algorithm
 - Viterbi algorithm



JOINT PROBABILITY $P(O, Q)$

- $P(Q|\lambda) = P(q_1) \cdot \prod_{t=2 \sim T} P(q_t|q_{t-1}, \lambda)$
 $= \pi_{q_1} \cdot a_{q_1 q_2} \cdot a_{q_2 q_3} \cdot \dots \cdot a_{q_{T-1} q_T} \quad (Q = q_1, q_2, \dots, q_T)$
- $P(O|Q, \lambda) = \prod_{t=1 \sim T} P(O_t|q_t, \lambda)$
 $= b_{q_1}(O_1) \cdot b_{q_2}(O_2) \cdot \dots \cdot b_{q_T}(O_T)$
- $P(O, Q|\lambda) = P(Q|\lambda) \cdot P(O|\lambda, Q)$
 $= \pi_{q_1} \cdot b_{q_1}(O_1) \cdot a_{q_1 q_2} \cdot b_{q_2}(O_2) \cdot \dots \cdot a_{q_{T-1} q_T} \cdot b_{q_T}(O_T)$



$$P(Q) = \pi_1 a_{11} a_{12} a_{23} a_{31} a_{13} a_{33} a_{32}$$

$$P(O|Q) = b_1(O_1) b_1(O_2) b_2(O_3) b_3(O_4) \cdot b_1(O_5) b_3(O_6) b_3(O_7) b_2(O_8)$$




ISSUE OF COMPUTATION COMPLEXITY

- $P(O|\lambda) = \sum_Q P(O, Q|\lambda)$
 $= \sum_Q P(Q|\lambda) \cdot P(O|\lambda, Q)$
 $= \sum_Q [\pi_{q_1} \cdot b_{q_1}(O_1) \cdot a_{q_1 q_2} \cdot b_{q_2}(O_2) \cdot \dots \cdot a_{q_{T-1} q_T} \cdot b_{q_T}(O_T)]$
- $P(O|\lambda)$ is the *summation of $P(O, Q|\lambda)$ over ALL paths* on the state space
 - Note: path Q is *hidden* in HMM (cannot be seen)
 - $P(X) = \sum_Y P(X, Y)$
- On $N \times T$ state space, there are totally *N^T paths*
 - Infeasible if every path is calculated *independently*
 - $(2T-1) \cdot N^T$ multiplications and N^T-1 additions



FORWARD PROCEDURE


- $\alpha_t(j) \equiv P(O_1, \dots, O_t, q_t = S_j | \lambda)$
- $\alpha_1(j) = P(O_1, q_1 = S_j | \lambda)$
 $= P(q_1 = S_j | \lambda) \cdot P(O_1 | \lambda, q_1 = S_j)$
 $= \pi_j \cdot b_j(O_1)$



- $\alpha_{t+1}(j) \equiv P(O_1, \dots, O_{t+1}, q_{t+1} = S_j | \lambda)$
 $= P(O_1, \dots, O_t, q_{t+1} = S_j | \lambda) \cdot P(O_{t+1} | \lambda, O_1, \dots, O_t, q_{t+1} = S_j)$
 $= [\sum_i P(O_1, \dots, O_t, q_t = S_i, q_{t+1} = S_j | \lambda)] \cdot b_j(O_{t+1})$
 $= \{\sum_i [P(O_1, \dots, O_t, q_t = S_j | \lambda) \cdot P(q_{t+1} = S_j | \lambda, O_1, \dots, O_t, q_t = S_i)]\} \cdot b_j(O_{t+1})$
 $= [\sum_i (\alpha_t(i) \cdot a_{ij})] \cdot b_j(O_{t+1})$



Stochastic Dependency

- $P(O | \lambda) = \sum_i \alpha_T(i)$



FORWARD PROCEDURE (CONT'D)

- Calculate $\alpha_t(i)$ for every *grid location* (t,i) on state space in left-to-right direction
- $\alpha_t(i)$ is the temporary summation of $P(O, Q|\lambda)$ over all partial paths terminating at (t, i)
- Dynamic Programming
- Complexity $\doteq N^2T$



BACKWARD PROCEDURE

- $\beta_t(i) \equiv P(O_{t+1}, \dots, O_T | q_t = S_i, \lambda)$

$$= \sum_j P(O_{t+1}, \dots, O_T, q_{t+1} = S_j | q_t = S_i, \lambda)$$


$$= \sum_j \{ P(q_{t+1} = S_j | q_t = S_i, \lambda) \cdot P(O_{t+1}, \dots, O_T | q_t = S_i, q_{t+1} = S_j, \lambda) \}$$

$$= \sum_j \{ a_{ij} \cdot [P(O_{t+1} | q_{t+1} = S_j, \lambda) \cdot P(O_{t+2}, \dots, O_T | q_{t+1} = S_j, O_{t+1}, \lambda)] \}$$

$$= \sum_j [a_{ij} \cdot b_j(O_{t+1}) \cdot \beta_{t+1}(j)]$$
- $\beta_{T-1}(i) = P(O_T | q_{T-1} = S_i, \lambda) = \sum_j P(O_T, q_T = S_j | q_{T-1} = S_i, \lambda)$

$$= \sum_j [P(q_T = S_j | q_{T-1} = S_i, \lambda) \cdot P(O_T | q_T = S_j, q_{T-1} = S_i, \lambda)]$$


$$= \sum_j [a_{ij} \cdot b_j(O_T)]$$
- It is there appropriate to assign $\beta_T(i) \equiv 1$ such that the inductive formula holds for $t = T-1$
- $\beta_t(i)$ is calculated for every grid point (t, i) in right-to-left direction



FIND OPTIMAL STATES Q_T^*

- $P(O, q_t = S_i | \lambda) = P(O_1, \dots, O_t, O_{t+1}, \dots, O_T, q_t = S_i | \lambda)$
 $= P(O_1, \dots, O_t, q_t = S_i | \lambda) \cdot P(O_{t+1}, \dots, O_T | q_t = S_i, O_1, \dots, O_t, \lambda)$
 $= \alpha_t(i) \cdot \beta_t(i)$
 - Summation of $P(O, Q | \lambda)$ over all paths passing (t, i)

- $q_t^* \equiv \operatorname{argmax}_i \{ P(q_t = S_i | O, \lambda) \}$
 $= \operatorname{argmax}_i \{ P(O, q_t = S_i | \lambda) / P(O | \lambda) \}$
 $= \operatorname{argmax}_i \{ \gamma_t(i) \}$
 - $\gamma_t(i) \equiv P(q_t = S_i | O, \lambda) = \alpha_t(i) \cdot \beta_t(i) / P(O | \lambda)$



- q_t^* 's are the states optimized for each time **individually**
- $P(q_1^*, q_2^*, \dots, q_T^*, O | \lambda)$ is **NOT** the path with highest probability $P(Q, O | \lambda)$ among all paths Q 's




FIND OPTIMAL PATH - VITERBI ALGORITHM

- $P^* = \max_Q P(O, Q | \lambda), Q^* = \operatorname{argmax}_Q P(O, Q | \lambda)$


Q^* is the path with the highest probability
 P^*

- $\delta_t(i) \equiv \max_{q_1, q_2, \dots, q_{t-1}} P(O_1, \dots, O_t, q_t = S_i | \lambda)$

- $\delta_t(i)$ is the temporary maximum of $P(O, Q | \lambda)$ for all partial paths terminating at (t, i)



- $\delta_t(j) = \pi_j \cdot b_j(O_1), \varphi_t(j) = 0$
- $\delta_{t+1}(j) = \max_i [\delta_t(i) \cdot a_{ij}] \cdot b_j(O_{t+1})$
 $\psi_{t+1}(j) = \operatorname{argmax}_i [\delta_t(i) \cdot a_{ij}]$
- $P^* = \max_i [\delta_T(i)]$
 $q_T' = \operatorname{argmax}_i [\delta_T(i)]$
 $q_t' = \varphi_{t+1}(q_{t+1}')$ backtracking
 $Q^* = q_1', q_2', \dots, q_T'$



At any time, only the **winner** can survive and be memorized

VITERBI ALGORITHM (CONT'D)

- Take *log*, the multiplications becomes additions


- $\delta'_{t+1}(j) = \max_i [\delta'_t(i) + \log(a_{ij})] + \log[b_j(O_{t+1})]$

- $\psi_{t+1}(j) = \operatorname{argmax}_i [\delta'_t(i) + \log(a_{ij})]$

- Avoid underflow for long observation sequence


- Concept analogy

- To pick up the maximum points on the street
 - Can proceed in directions \rightarrow and \nearrow
 - Get $\log(a_{ij})$ points by transit from i to j
 - Get $\log[b_j(O_t)]$ points on the cross (t, j)



VITERBI ALGORITHM (CONT'D)

- For **left-to-right** HMM, the optimal path Q^* obtained from Viterbi algorithm can **segment** the observation sequence into different states
- Can be used for **automatic segmentation**




EM ALGORITHM

- Expectation-Maximization (8.30)
- Find λ^* by iterative procedure
 - $P(O | \lambda') > P(O | \lambda)$ and can find **local maximum**
 - Converge** when the change of likelihood is small
- $\gamma_t(i) \equiv P(q_t = S_i | O, \lambda)$ *the contribution of O_t for state i*
- $\varepsilon_t(i, j) \equiv P(q_t = S_i, q_{t+1} = S_j | O, \lambda)$

$$\gamma_t(i) = \frac{\alpha_t(i) \beta_t(i)}{P(O | \lambda)}$$

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{P(O | \lambda)}$$



EM ALGORITHM (CONT'D)

◦ Baum-Welch Reestimation Procedure

$\bar{\pi}_i$ = expected frequency (number of times) in state S_i at time $(t = 1) = \gamma_1(i)$

\bar{a}_{ij} = $\frac{\text{expected number of transitions from state } S_i \text{ to state } S_j}{\text{expected number of transitions from state } S_i}$

$$= \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

Σ_t : summation the contributions of all observations $\{O_t\}$

$\bar{b}_j(k)$ = $\frac{\text{expected number of times in state } j \text{ and observing symbol } v_k}{\text{expected number of times in state } j}$

$$= \frac{\sum_{t=1}^T \gamma_t(j) \text{ s.t. } O_t = v_k}{\sum_{t=1}^T \gamma_t(j)}$$

For Discrete state observation distribution

$$b_j(O_t = \underline{v}_k)$$



EM ALGORITHM (CONT'D)

For Continuous state observation distribution

$$b_j(O_t=x) = \sum_k c_{jk} \cdot N(x; \mu_{jk}, U_{jk})$$

$$\bar{c}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k)}{\sum_{t=1}^T \sum_{k=1}^M \gamma_t(j, k)}$$

$$\bar{\mu}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k) \cdot \mathbf{O}_t}{\sum_{t=1}^T \gamma_t(j, k)}$$

$$\bar{U}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k) \cdot (\mathbf{O}_t - \mu_{jk})(\mathbf{O}_t - \mu_{jk})'}{\sum_{t=1}^T \gamma_t(j, k)}$$

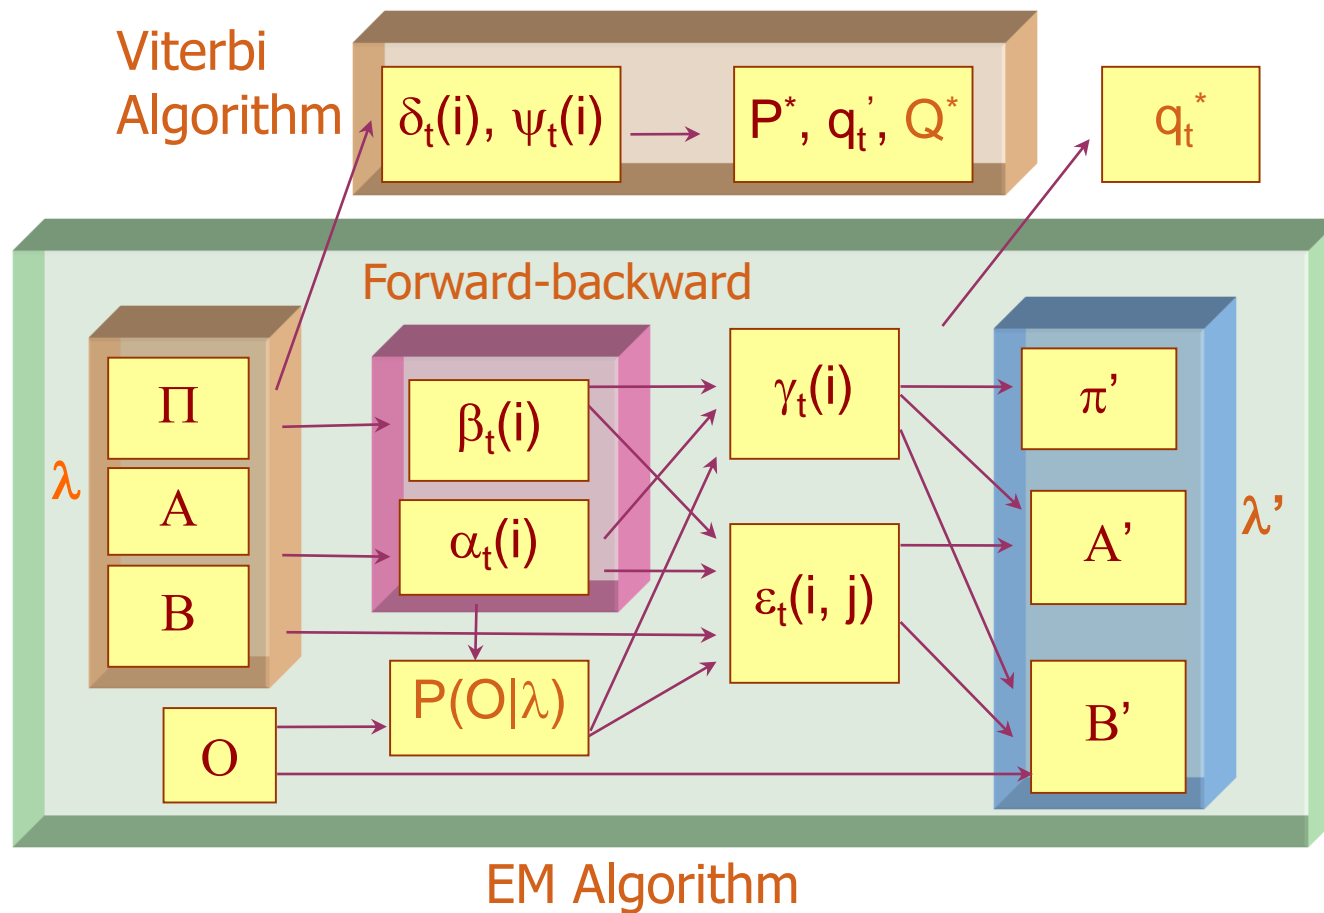
$$\gamma_t(j, k) = \left[\frac{\alpha_t(j) \beta_t(j)}{\sum_{j=1}^N \alpha_t(j) \beta_t(j)} \right] \left[\frac{c_{jk} \mathcal{N}(\mathbf{O}_t, \mu_{jk}, U_{jk})}{\sum_{m=1}^M c_{jm} \mathcal{N}(\mathbf{O}_t, \mu_{jm}, U_{jm})} \right]$$

$\gamma_t(j)$: How likely O_t belongs to state j

How likely O_t belongs to mixture k if in state j



DEPENDENCY DIAGRAM



HMM FOR PATTERN RECOGNITION

- Assume there are M patterns to be recognized, and we have observations for these patterns. Then, we can train the model λ_i for i -th pattern using the observations of this pattern
- $i^* = \operatorname{argmax}_i P(O | \lambda_i) \quad i=1, 2, \dots, M$



TOPOLOGY OF HMM

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix}.$$

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} & 0 \\ 0 & a_{22} & a_{23} & a_{24} \\ 0 & 0 & a_{33} & a_{34} \\ 0 & 0 & 0 & a_{44} \end{bmatrix}.$$






Fig. 7. Illustration of 3 distinct types of HMMs. (a) A 4-state ergodic model. (b) A 4-state left-right model. (c) A 6-state parallel path left-right model.

SCALING ISSUE

- If the length of observation sequence is long, there may arise the problem of underflow since all probabilities are less than 1.
 - $P(O | \lambda) = \sum_Q \pi_{q_1} \cdot b_{q_1}(O_1) \cdot a_{q_1 q_2} \cdot b_{q_2}(O_2) \cdot \dots \cdot a_{q_{T-1} q_T} \cdot b_{q_T}(O_T)$
- $S_t = 1 / \sum_i \alpha_t(i)$ is used for scaling $\alpha_t(i)$ and $\beta_t(i)$ at time t in Forward-backward procedure
 - It can be proved scaling will NOT influence the reestimation formula
- In Viterbi algorithm, probability is taken “log” and then accumulated



MULTIPLE OBSERVATION SEQUENCES

$$\bar{a}_{ij} = \frac{\sum_{k=1}^K \frac{1}{P_k} \sum_{t=1}^{T_k-1} \alpha_t^k(i) a_{ij} b_j(O_{t+1}^{(k)}) \beta_{t+1}^k(j)}{\sum_{k=1}^K \frac{1}{P_k} \sum_{t=1}^{T_k-1} \alpha_t^k(i) \beta_t^k(j)}$$

$$\bar{b}_j(i) = \frac{\sum_{k=1}^K \frac{1}{P_k} \sum_{\substack{t=1 \\ \text{s.t. } O_t = v_i}}^{T_k-1} \alpha_t^k(i) \beta_t^k(j)}{\sum_{k=1}^K \frac{1}{P_k} \sum_{t=1}^{T_k-1} \alpha_t^k(i) \beta_t^k(j)}$$

- The Reestimation formula can be modified for multiple observation sequences
- $P_k = P(O^{(k)} | \lambda)$



INITIAL MODEL

- Good initial estimate is important for **avoiding local maximum**
- Random or uniform distributed parameters are adequate for Π and A , but not enough for B .
- To obtain better initial estimate
 - Manual segmentation
 - Segmental K-means (Viterbi segmentation)




LIMITATION OF HMM

- First-Order Assumption
 - Markov Chain
 - Leading to exponential duration model
- Conditional Independence Assumption
 - The Observation depends **ONLY** on the **CURRENT** state
 - Maximum entropy Markov model (MEMM)
 - Conditional random field (CRF)



LIMITATION OF HMM (CONT'D)

- Local optimum guaranteed by EM algorithm




GENERALIZATION OF MM/HMM

- Observation O can be **vector** instead of scalar
- Concept of **time** can be generalized as other physical dimension
 - Example: x-axis in Euclidean space
 - State sequence $Q = \{ q_t \}$
 - Observation sequence $O = \{ O_t \}$
 - Recognition of patterns which vary with some dimension (e.g. image, gene pattern, ...)



GENERALIZATION OF MM/HMM

- State can be flexibly defined
 - Number of state might be very large
- Better duration model
 - Use Gaussian or Gamma function as duration distribution instead of exponential function (a_{ij} is exponential duration distribution effectively)
 - Slight improvement achieved



APPLICATIONS OF HMM

- Statistical Model for *Trajectory*
- Speech recognition/synthesis
 - Spectrum trajectory
- Image recognition
 - row trajectory
- Natural language processing
 - Word segmentation / POS tagging
- Gesture recognition
 - 2D/3D trajectory
 - Action recognition (multi-point trajectory)
- Melody recognition/synthesis
 - Note trajecotry

